

# Neural Networks

Andreas Müller

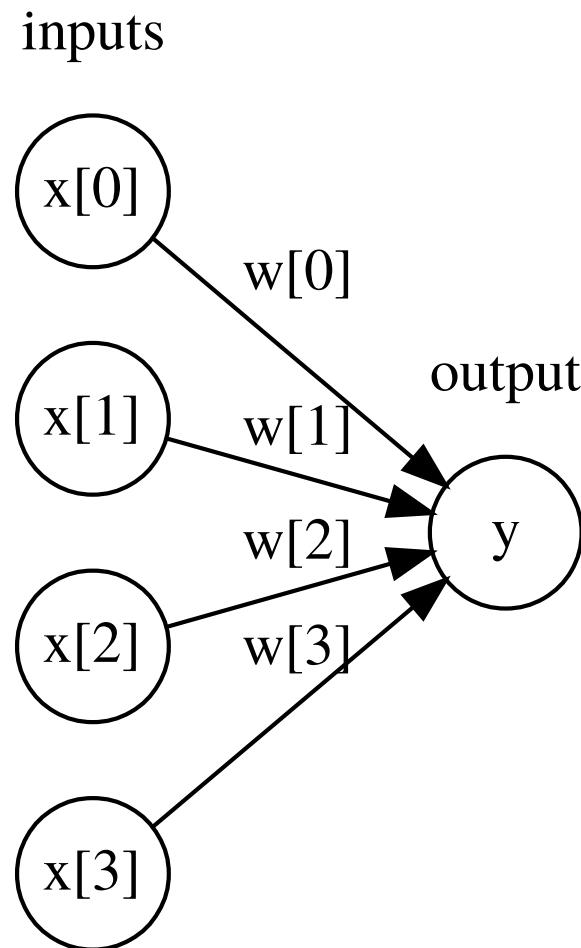
# History

- Nearly everything we talk about today existed ~1990
- What changed?
  - More data
  - Faster computers (GPUs)
  - Some improvements:
    - relu
    - Drop-out
    - adam
    - Batch-normalization
    - Residual networks

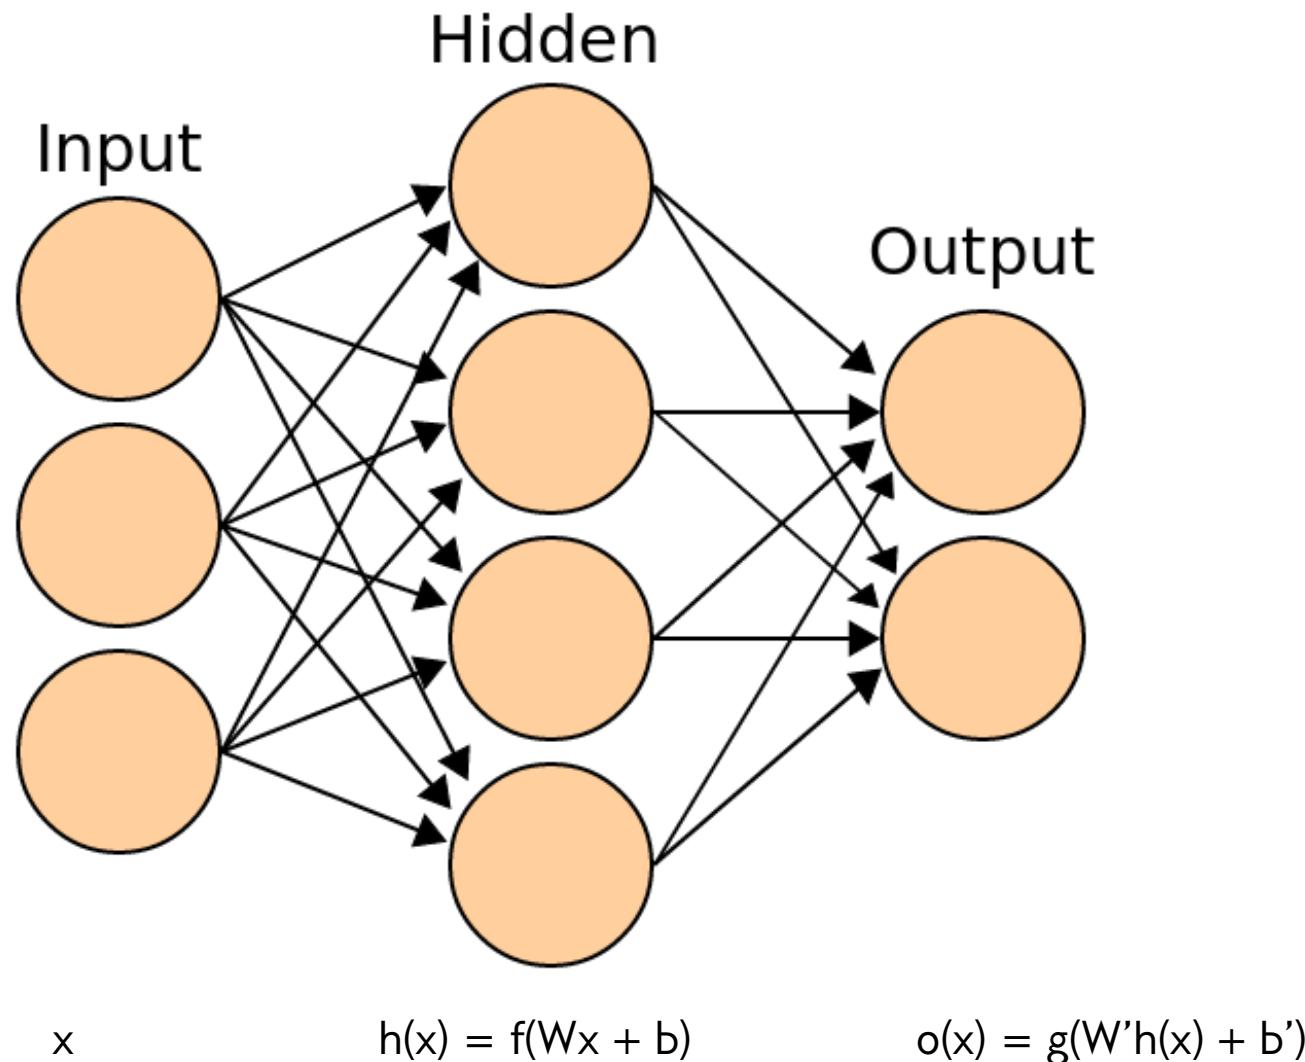
# Supervised Neural Networks

- Non-linear models for classification and regression
- Work well for very large datasets
- Non-convex optimization
- Notoriously slow to train – need for GPUs
- Use dot products etc → require preprocessing, similar to SVM or linear models, unlike trees
- MANY variants (Convolutional nets, Gated Recurrent neural networks, Long-Short-Term Memory, recursive neural networks, variational autoencoders, general adversarial networks, neural turing machines...)

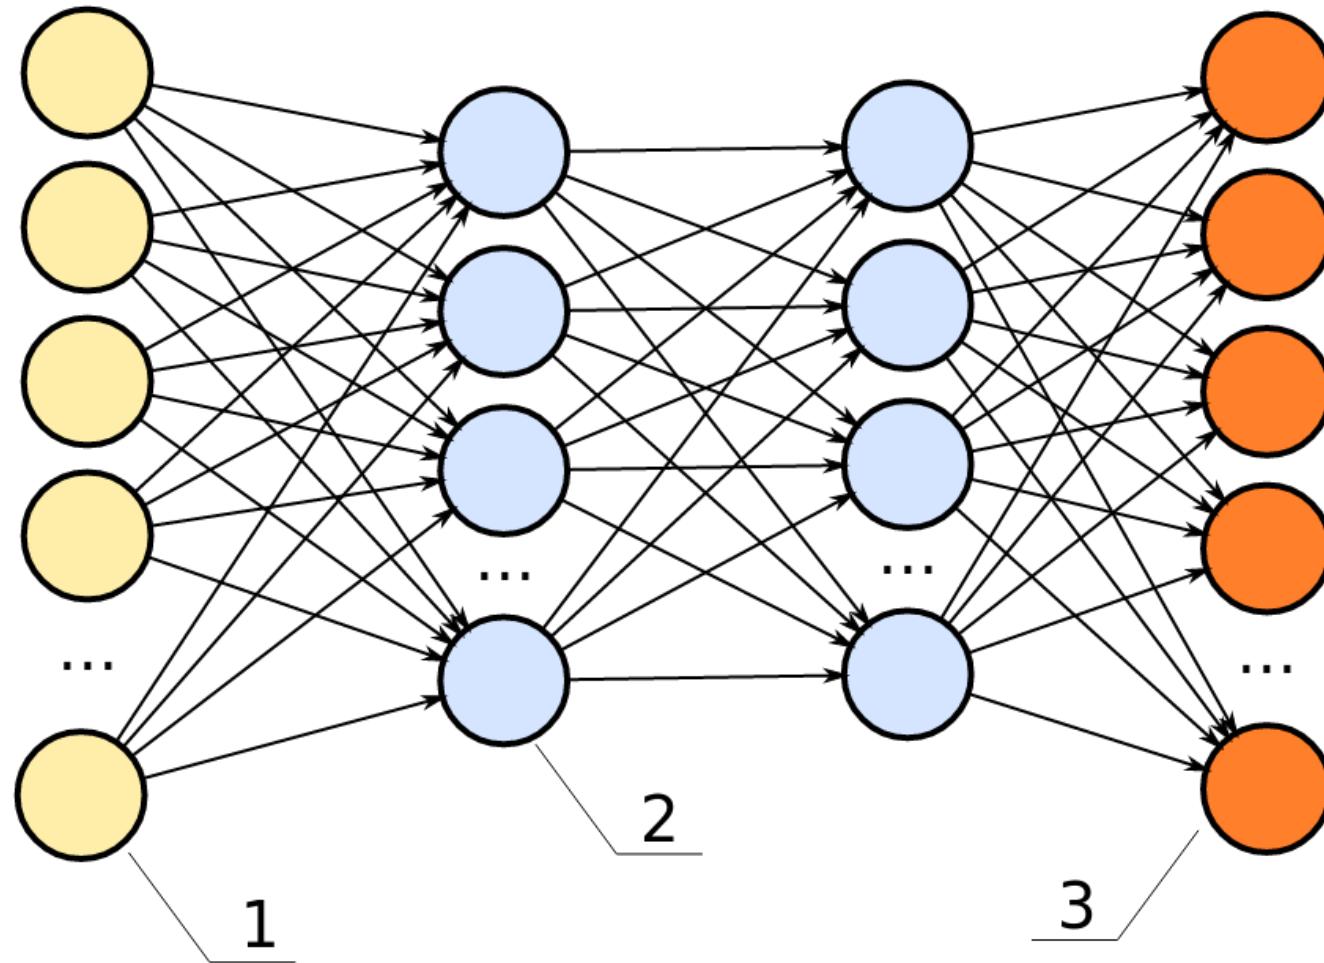
# Logistic regression drawn as neural net



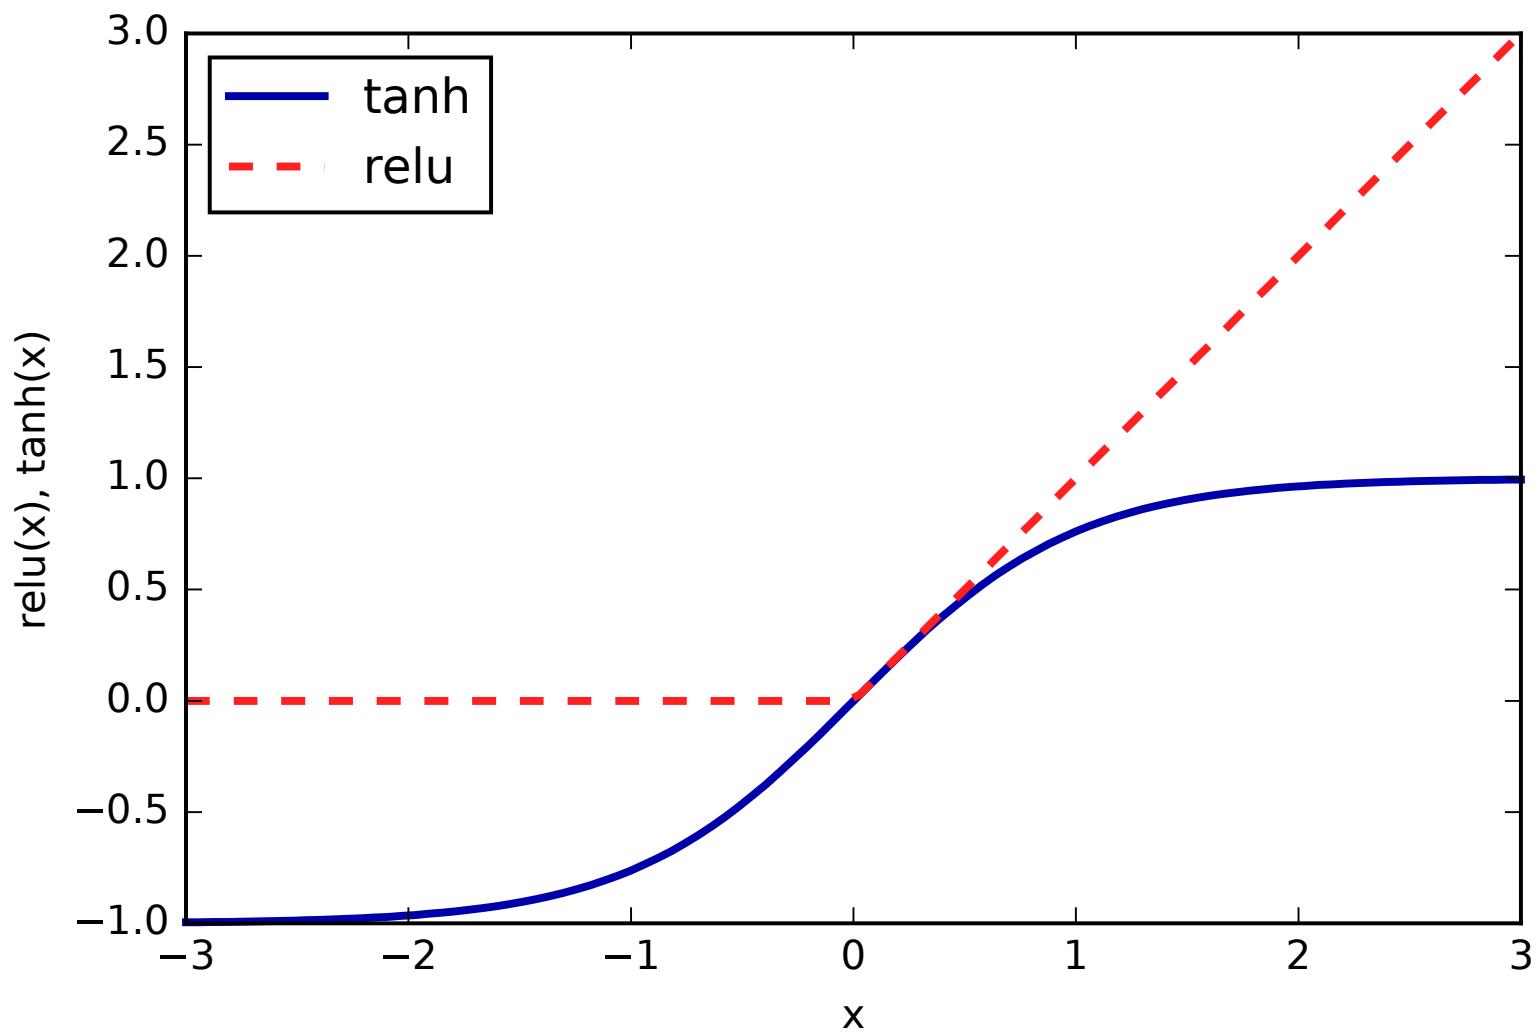
# Basic Architecture (for making predictions)



# Can have arbitrary many layers



# Nonlinear activation function



# Training objective

$$h(\mathbf{x}) = f(W_1 \mathbf{x} + \mathbf{b}_1)$$

$$o(\mathbf{x}) = g(W_2 h(\mathbf{x}) + \mathbf{b}_2) = g(W_2 f(W_1 \mathbf{x} + \mathbf{b}_1) + \mathbf{b}_2)$$

$$\min_{W_1, W_2, \mathbf{b}_1, \mathbf{b}_2} \sum_{i=1}^N \ell(y_i, o(\mathbf{x}_i)) \quad \text{Could add regularization}$$

$$= \min_{W_1, W_2, \mathbf{b}_1, \mathbf{b}_2} \sum_{i=1}^N \ell(y_i, g(W_2 f(W_1 \mathbf{x}_i + \mathbf{b}_1) + \mathbf{b}_2))$$

$\ell$  squared loss for regression  
cross-entropy loss (multi-class log-loss) for classification

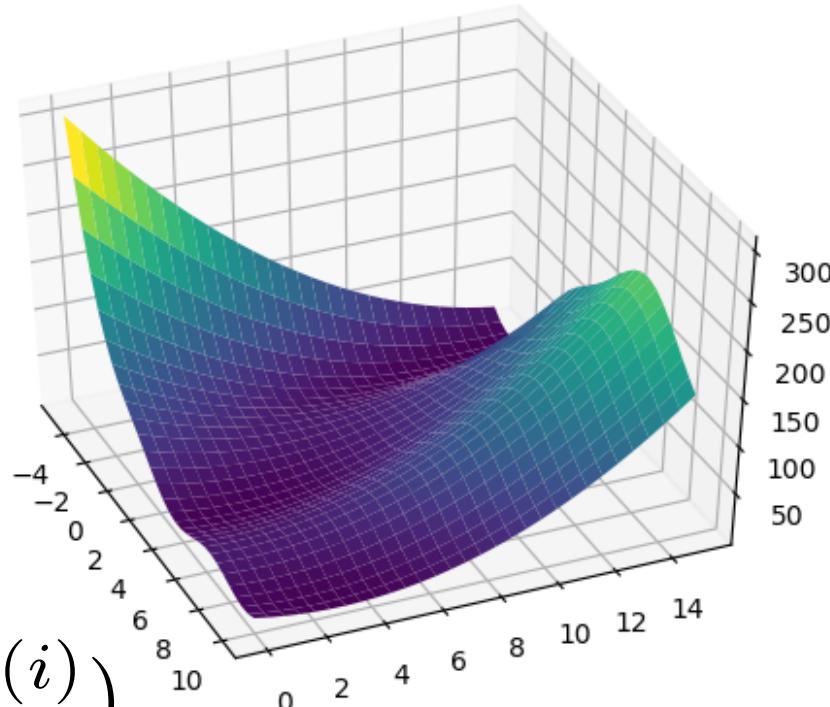
# Reminder: Gradient Descent

Want:  $\arg \min_w F(w)$

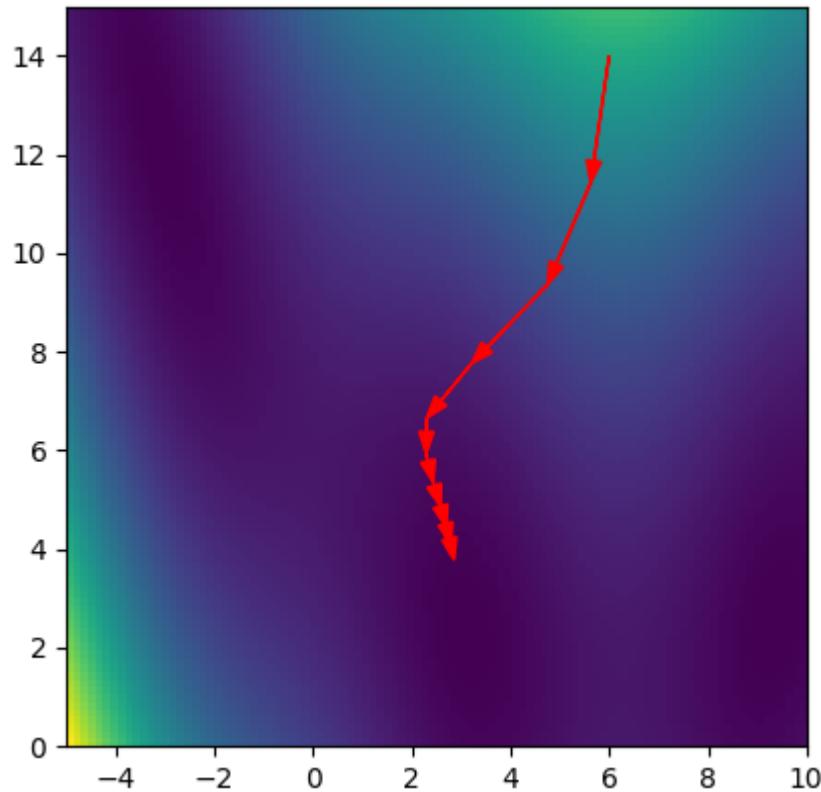
Initialize  $w_0$

$$w^{(i+1)} \leftarrow w^{(i)} - \eta_i \frac{d}{dw} F(w^{(i)})$$

Converges to local minimum

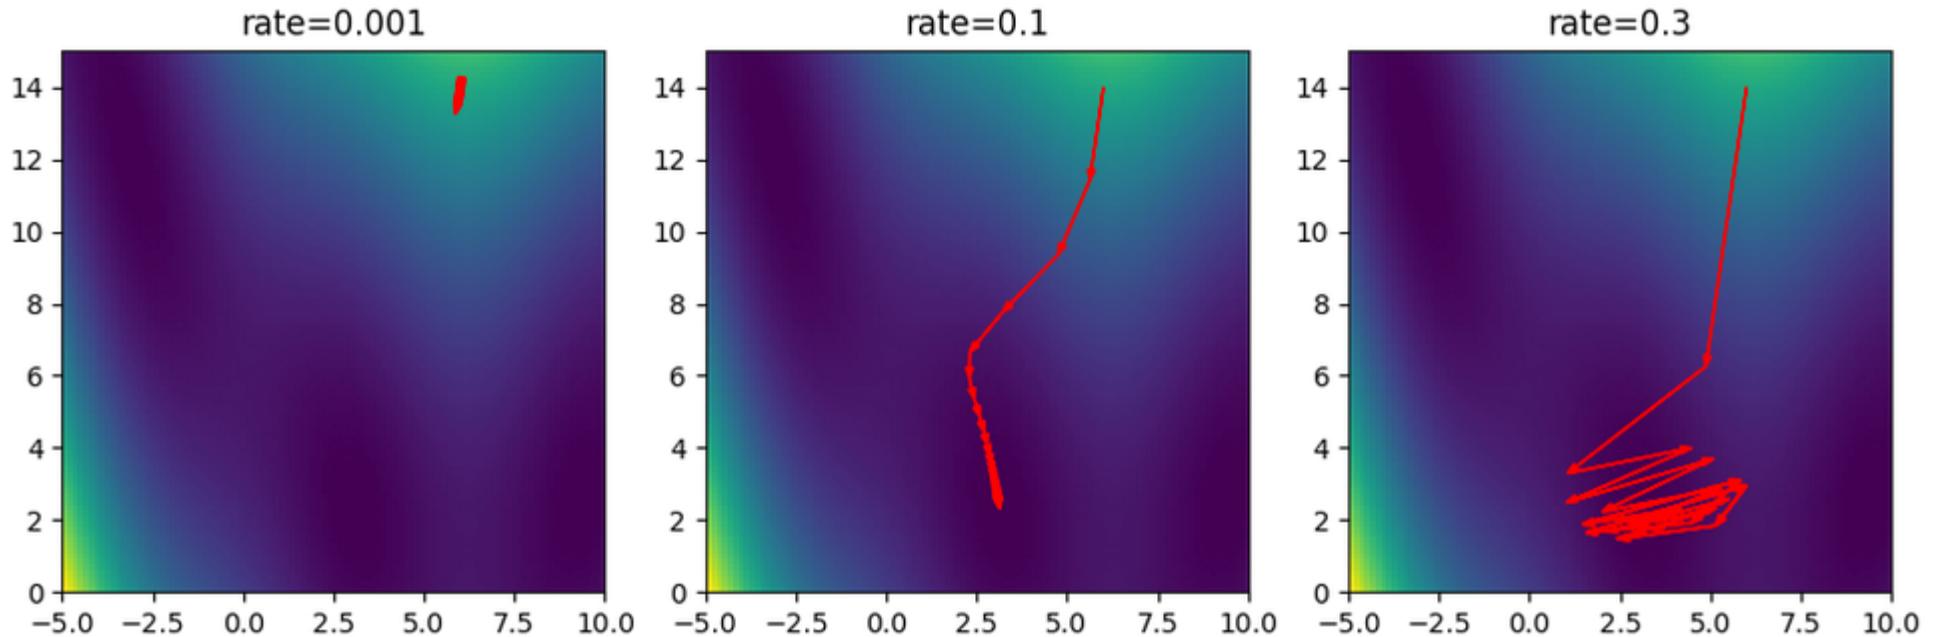


# Reminder: Gradient Descent



$$w^{(i+1)} \leftarrow w^{(i)} - \eta_i \frac{d}{dw} F(w^{(i)})$$

# Picking a learning rate



$$w^{(i+1)} \leftarrow w^{(i)} - \eta_i \frac{d}{dw} F(w^{(i)})$$

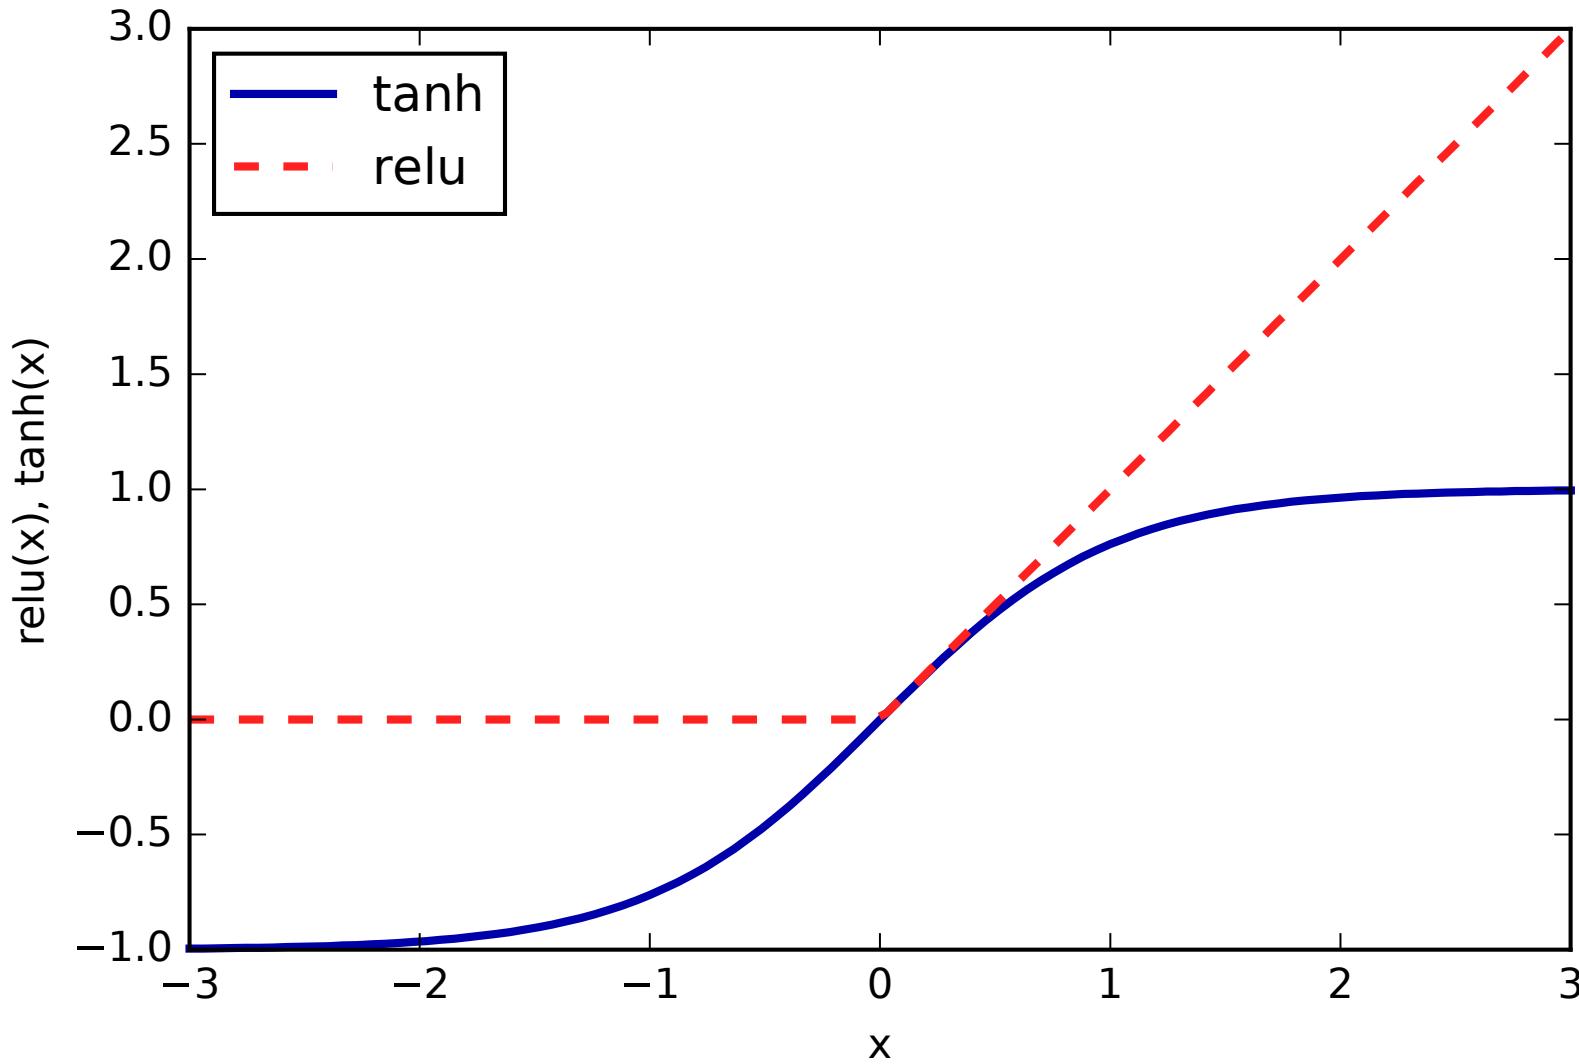
# Backpropagation

- For gradient based-method need  $\frac{\partial o(\mathbf{x})}{\partial W_i}$   $\frac{\partial o(\mathbf{x})}{\partial \mathbf{b}_i}$   
 $\text{net}(\mathbf{x}) := W_1 \mathbf{x} + b_1$

$$\frac{\partial o(\mathbf{x})}{\partial W_1} = \underbrace{\frac{\partial o(\mathbf{x})}{\partial h(\mathbf{x})}}_{\text{backpropagation of gradient of layer above.}} \underbrace{\frac{\partial h(\mathbf{x})}{\partial \text{net}(\mathbf{x})}}_{\text{Gradient of Non-linearity } f} \underbrace{\frac{\partial \text{net}(\mathbf{x})}{\partial W_1}}_{\text{Input to 1st layer } x}$$

Backpropagation = Chain Rule + Dynamic Programming

# But wait!



# Optimizing W, b

$$W_i \leftarrow W_i - \eta \sum_{j=1}^n \frac{\ell(\mathbf{x}_j, y_j)}{W_i} \quad \text{batch}$$

$$W_i \leftarrow W_i - \eta \sum_{j=k}^{k+m} \frac{\ell(\mathbf{x}_j, y_j)}{W_i} \quad \text{minibatch}$$

$$W_i \leftarrow W_i - \eta \frac{\ell(\mathbf{x}_j, y_j)}{W_i} \quad \text{Online / stochastic}$$

# Learning Heuristics

- Constant  $\eta$  not good
- Can decrease  $\eta$
- Better: adaptive  $\eta$  for each entry if  $W_i$
- State-of-the-art: adam (with magic numbers)

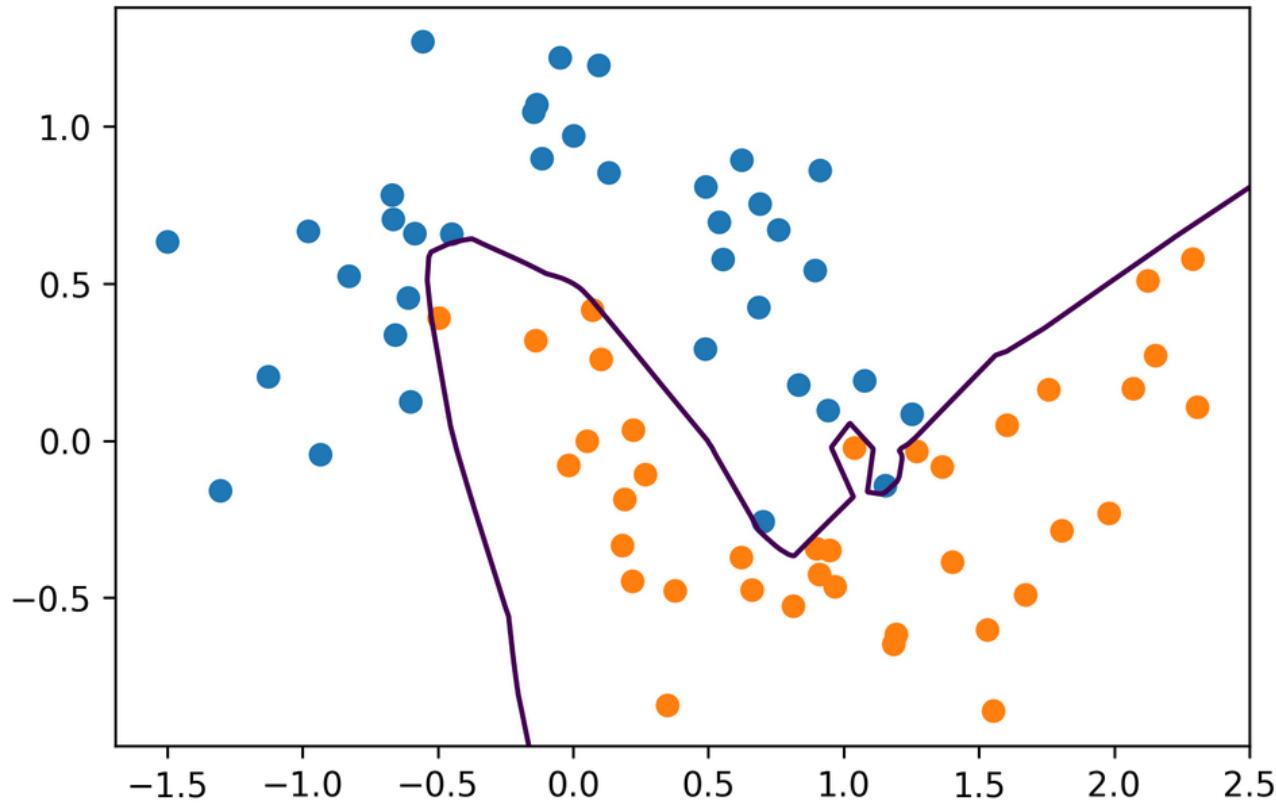
<https://arxiv.org/pdf/1412.6980.pdf>

<http://sebastianruder.com/optimizing-gradient-descent/>

# Picking optimization algorithms

- Small dataset: off the shelf like l-bfgs
- Big dataset: adam
- Have time & nerve: tune the schedule

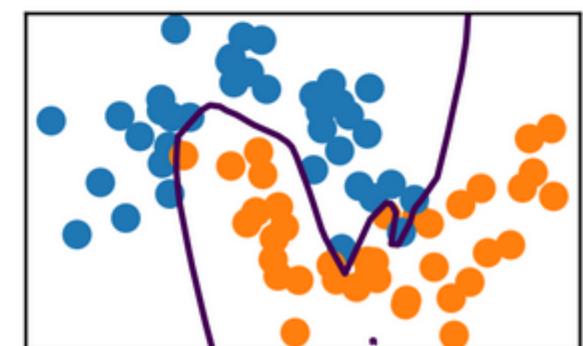
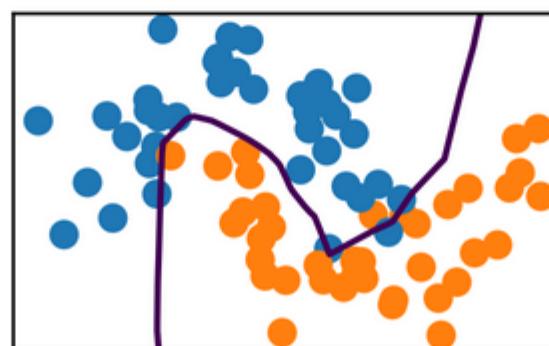
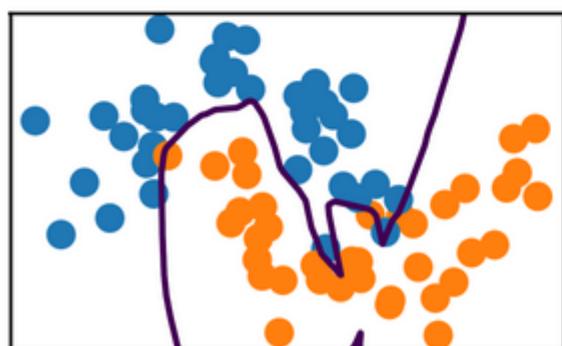
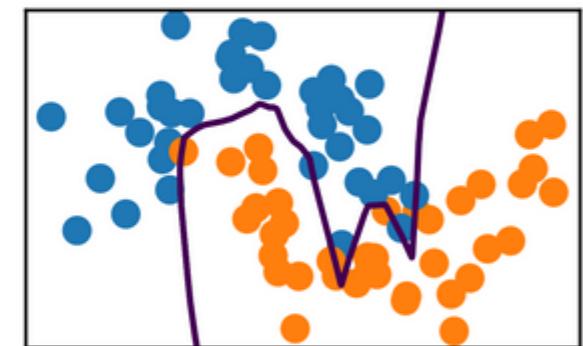
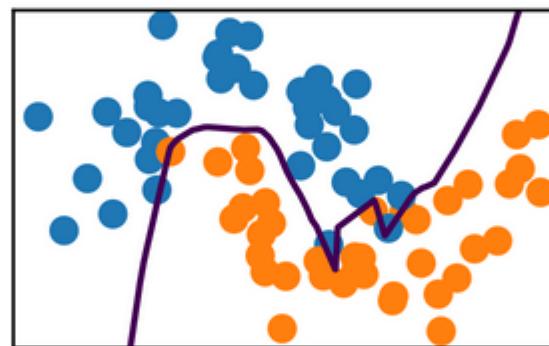
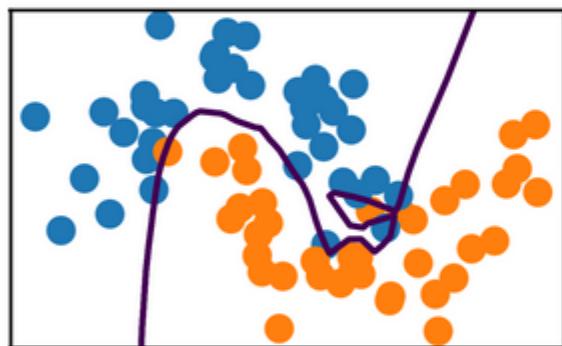
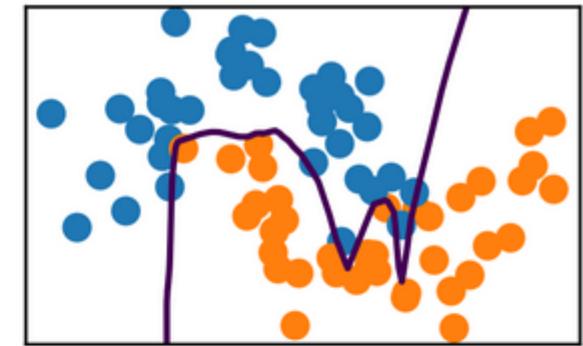
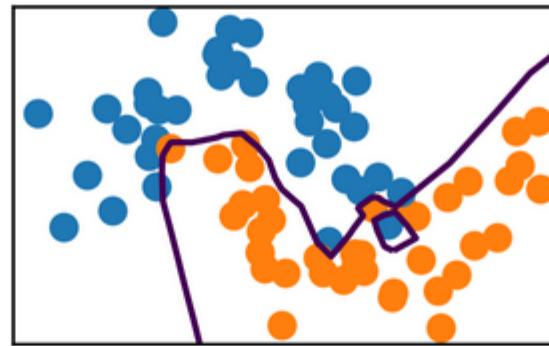
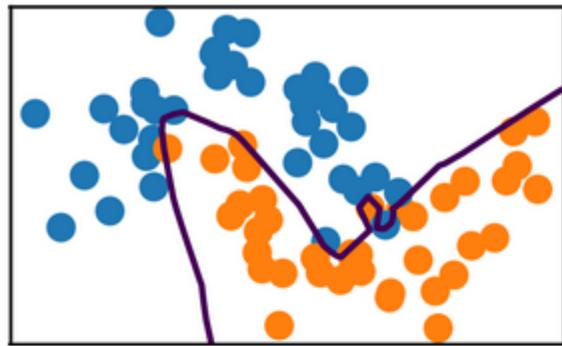
# Neural Nets with sklearn



```
mlp = MLPClassifier(solver='lbfgs', random_state=0).fit(X_train, y_train)
print(mlp.score(X_train, y_train))
print(mlp.score(X_test, y_test))
```

1.0  
0.88

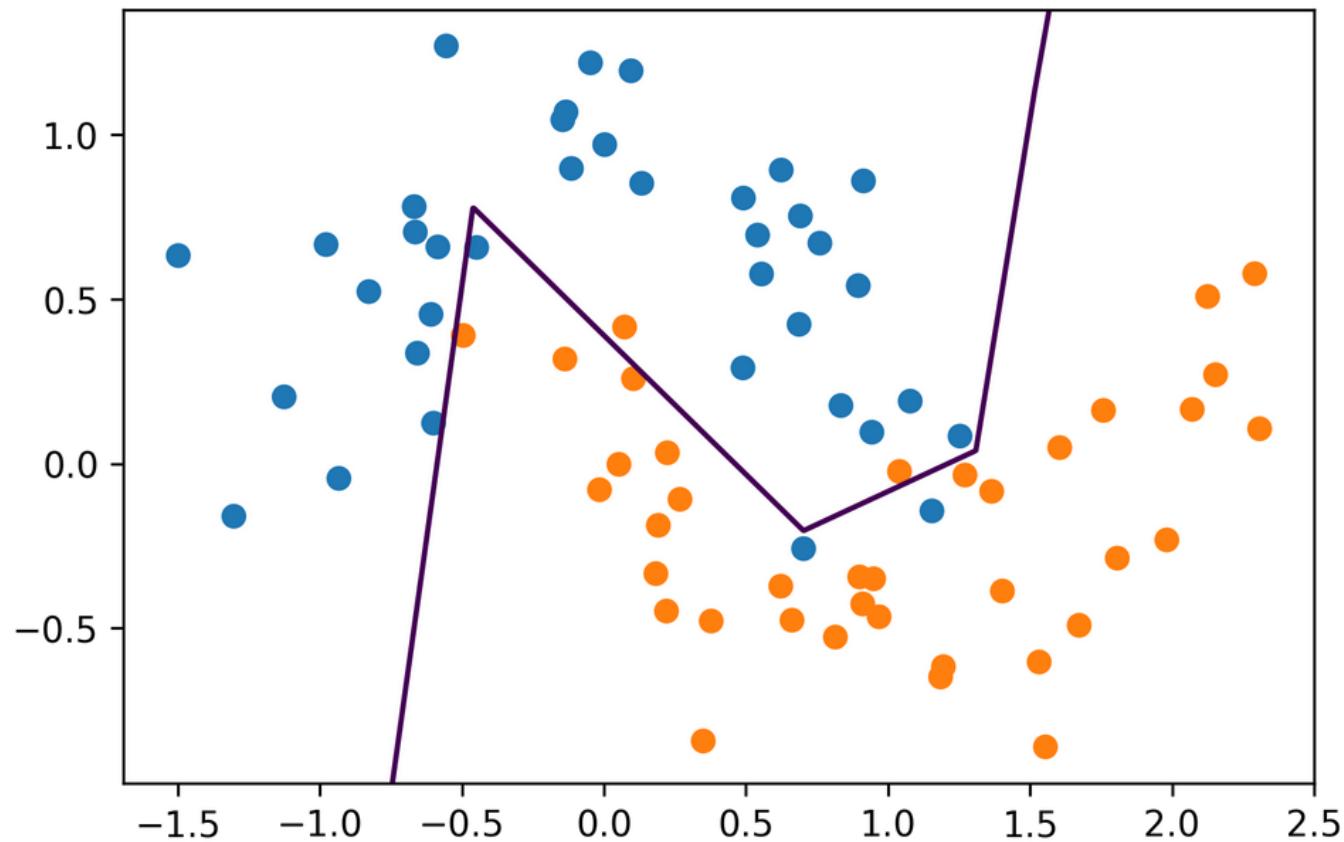
# Random State



# Hidden Layer Size

```
mlp = MLPClassifier(solver='lbfgs', hidden_layer_sizes=(5,), random_state=10)
mlp.fit(X_train, y_train)
print(mlp.score(X_train, y_train))
print(mlp.score(X_test, y_test))
```

0.933333333333  
0.92

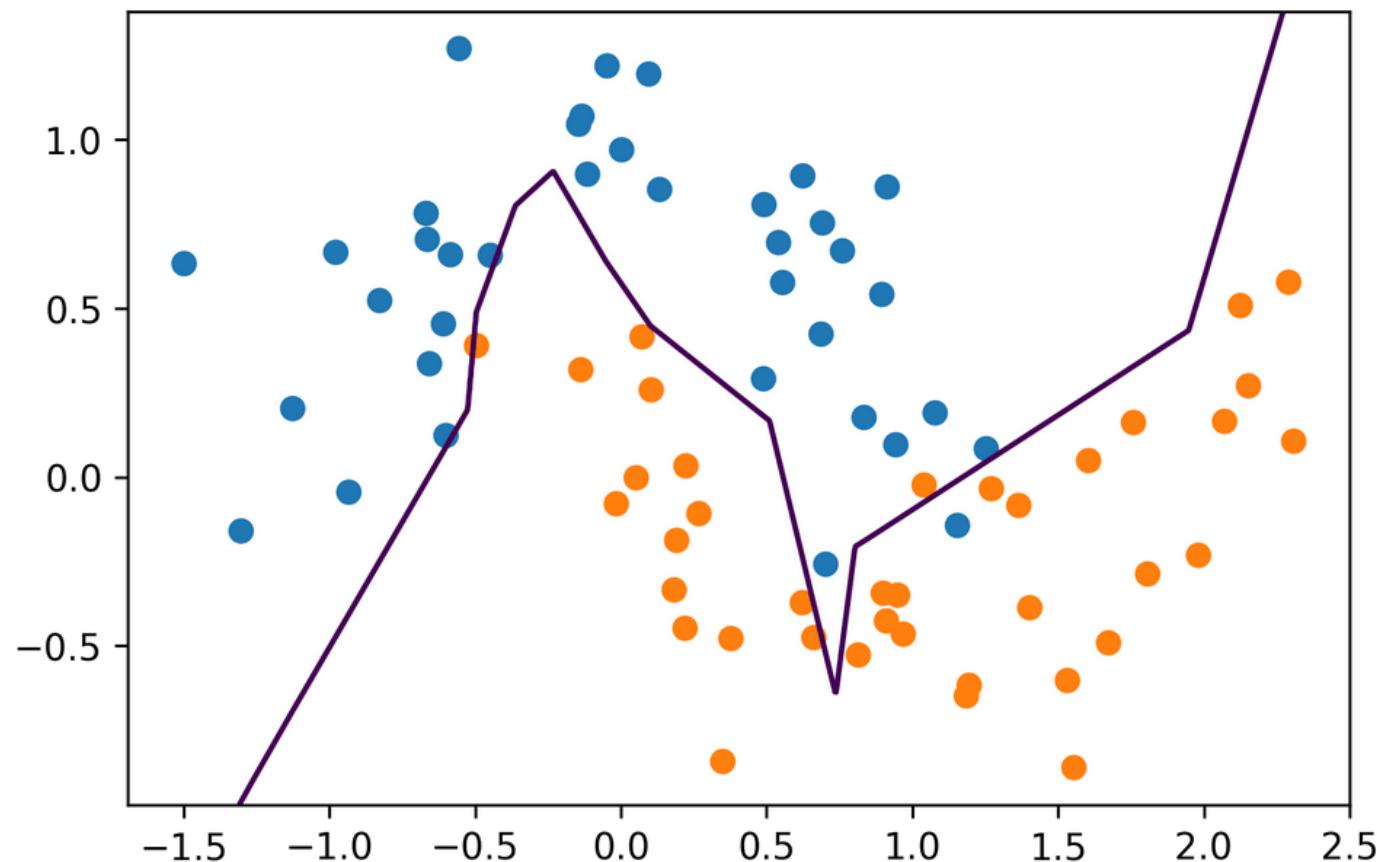


# Hidden Layer Size

```
mlp = MLPClassifier(solver='lbfgs', hidden_layer_sizes=(10, 10, 10), random_state=0)
mlp.fit(X_train, y_train)
print(mlp.score(X_train, y_train))
print(mlp.score(X_test, y_test))
```

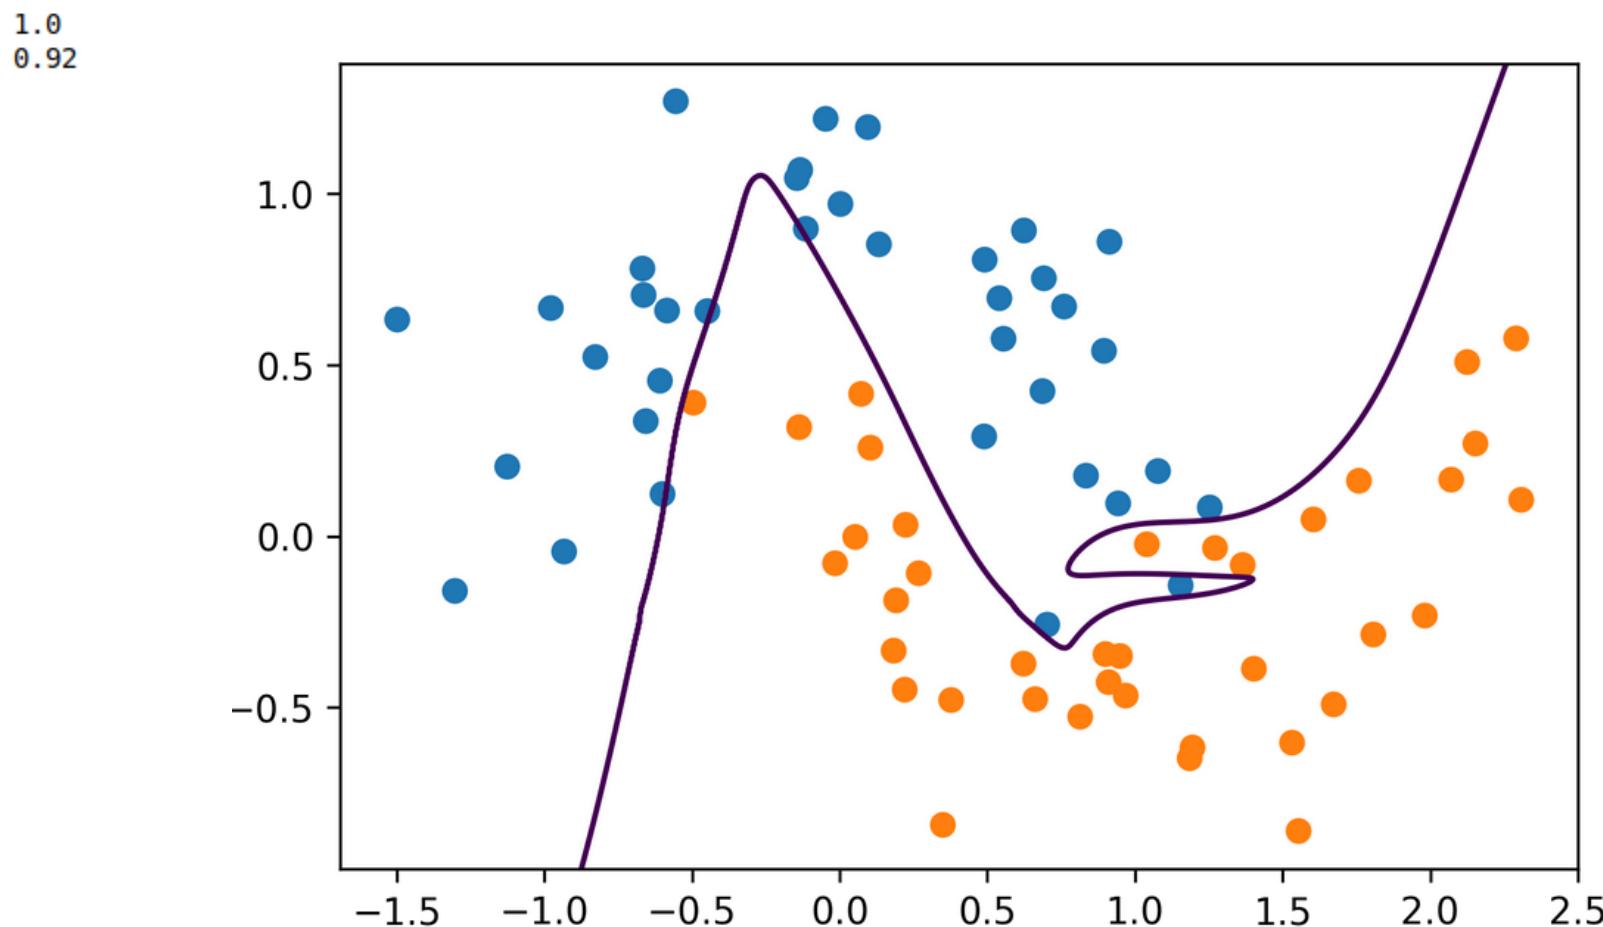
0.973333333333

0.84

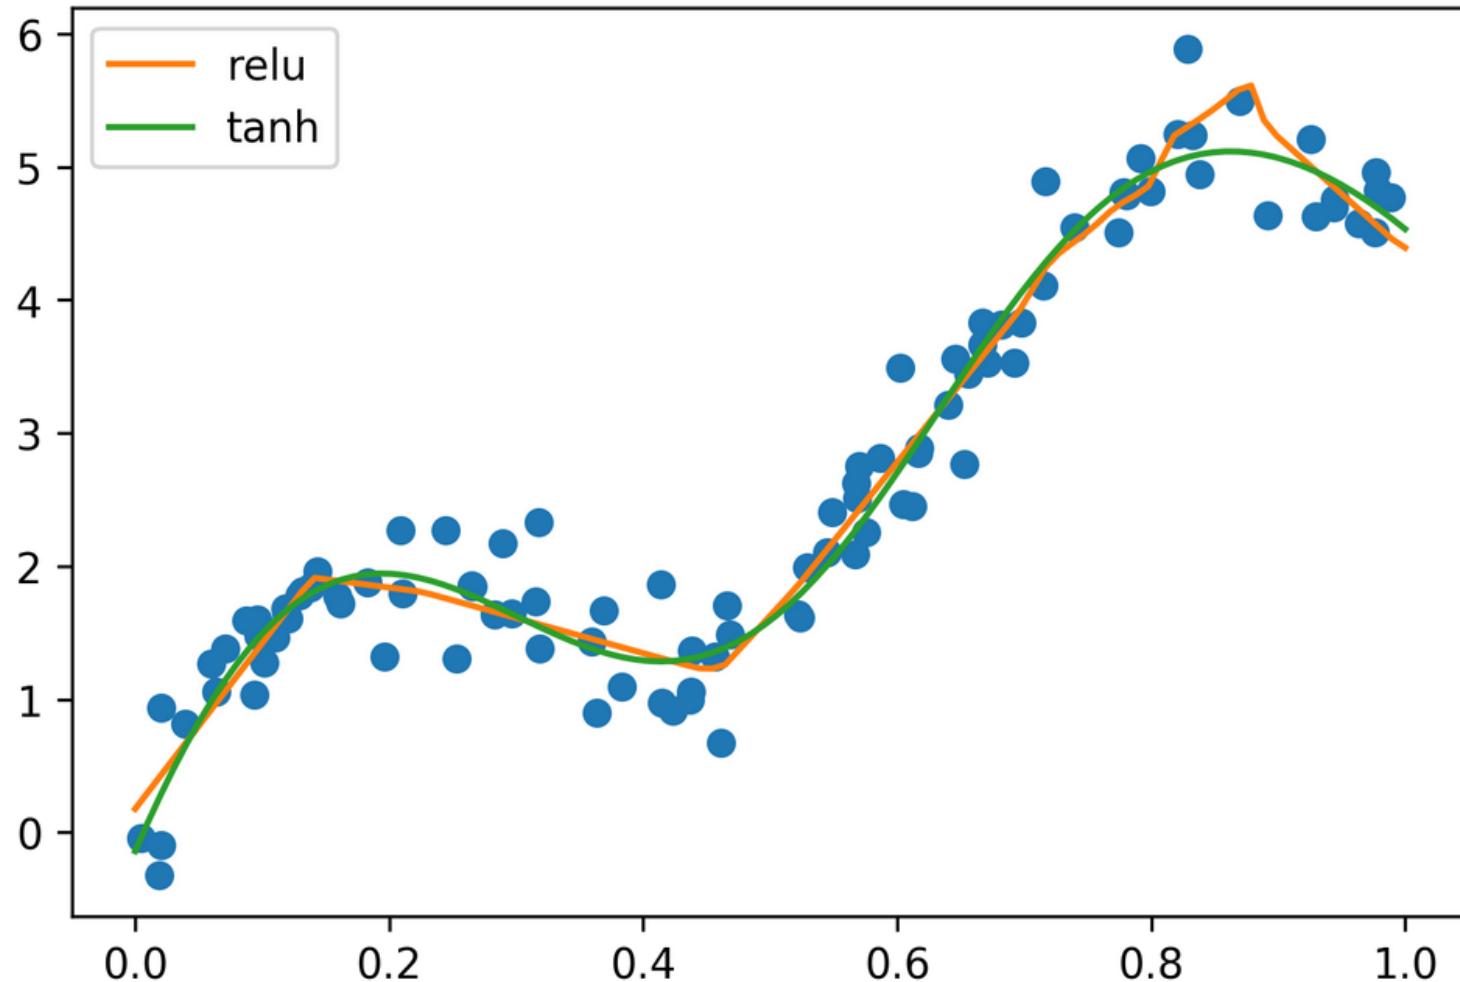


# Activation Functions

```
mlp = MLPClassifier(solver='lbfgs', hidden_layer_sizes=(10, 10, 10), activation="tanh", random_state=0)
mlp.fit(X_train, y_train)
print(mlp.score(X_train, y_train))
print(mlp.score(X_test, y_test))
```



# Regression



```
from sklearn.neural_network import MLPRegressor  
mlp_relu = MLPRegressor(solver="lbfgs").fit(X, y)  
mlp_tanh = MLPRegressor(solver="lbfgs", activation='tanh').fit(X, y)
```

# Complexity Control

- Number of parameters
- Regularization
- Early stopping
- (drop-out)

Getting flexible & Scaling up

# Deep Learning Libraries

- tf.learn (Tensorflow)
- Keras (Tensorflow, Theano)
- Lasagna (Theano)
- Torch.nn / PyTorch (torch)
- Chainer (chainer)
- MXNet (MXNet)
- Also see:  
[http://mxnet.io/architecture/program\\_model.html](http://mxnet.io/architecture/program_model.html)

# Quick look at TensorFlow

- “down to the metal” - don’t use for everyday tasks
- Three steps for learning:
  - Build the computation graph (using array operations and functions etc)
  - Create an Optimizer (gradient descent, adam, ...) attached to the graph.
  - Run the actual computation.

```
import tensorflow as tf
import numpy as np

# Create 100 phony x, y data points in NumPy, y = x * 0.1 + 0.3
x_data = np.random.rand(100).astype(np.float32)
y_data = x_data * 0.1 + 0.3

# create graph: model
W = tf.Variable(tf.random_uniform([1], -1.0, 1.0))
b = tf.Variable(tf.zeros([1]))
y = W * x_data + b

# create graph: loss
loss = tf.reduce_mean(tf.square(y - y_data))

# bind optimizer
optimizer = tf.train.GradientDescentOptimizer(0.5)
train = optimizer.minimize(loss)

# run graph
init = tf.global_variables_initializer()
sess = tf.Session()
sess.run(init)

# Fit the line.
for step in range(201):
    sess.run(train)
    if step % 20 == 0:
        print(step, sess.run(W), sess.run(b))
```

No computation

Allocate variables

All the work / computation

# Introduction to Keras

# Keras Sequential

```
from keras.models import Sequential
from keras.layers import Dense, Activation

model = Sequential([
    Dense(32, input_shape=(784,)),
    Activation('relu'),
    Dense(10),
    Activation('softmax'),
])
```

Using TensorFlow backend.

```
model = Sequential()
model.add(Dense(32, input_dim=784))
model.add(Activation('relu'))
```

```
model = Sequential([
    Dense(32, input_shape=(784,), activation='relu'),
    Dense(10, activation='softmax'),
])
```

```
model.summary()
```

Layer (type)	Output Shape	Param #
dense_165 (Dense)	(None, 32)	25120
activation_113 (Activation)	(None, 32)	0
dense_166 (Dense)	(None, 10)	330
activation_114 (Activation)	(None, 10)	0

Total params: 25,450.0

Trainable params: 25,450.0

Non-trainable params: 0.0

# Setting optimizer

## compile

```
compile(self, optimizer, loss, metrics=None, sample_weight_mode=None)
```

Configures the learning process.

## Arguments

- **optimizer**: str (name of optimizer) or optimizer object. See [optimizers](#).
- **loss**: str (name of objective function) or objective function. See [objectives](#).
- **metrics**: list of metrics to be evaluated by the model during training and testing. Typically you will use `metrics=['accuracy']`. See [metrics](#).

```
model.compile("adam", "categorical_crossentropy", metrics=['accuracy'])
```

# Training the model

```
fit(self, x, y, batch_size=32, epochs=10, verbose=1, callbacks=None, validation_split=0.0, validation_data=None, shuffle=True, initial_epoch=0)
```

Trains the model for a fixed number of epochs.

## Arguments

- **x**: input data, as a Numpy array or list of Numpy arrays (if the model has multiple inputs).
- **y**: labels, as a Numpy array.
- **batch\_size**: integer. Number of samples per gradient update.
- **epochs**: integer, the number of epochs to train the model.
- **verbose**: 0 for no logging to stdout, 1 for progress bar logging, 2 for one log line per epoch.
- **callbacks**: list of `keras.callbacks.Callback` instances. List of callbacks to apply during training. See [callbacks](#).
- **validation\_split**: float ( $0 < x < 1$ ). Fraction of the data to use as held-out validation data.
- **validation\_data**: tuple  $(x_{\text{val}}, y_{\text{val}})$  or tuple  $(x_{\text{val}}, y_{\text{val}}, \text{val\_sample\_weights})$  to be used as held-out validation data. Will override `validation_split`.
- **shuffle**: boolean or str (for 'batch'). Whether to shuffle the samples at each epoch. 'batch' is a special option for dealing with the limitations of HDF5 data; it shuffles in batch-sized chunks.

# Preparing MNIST data

```
from keras.datasets import mnist
import keras

(X_train, y_train), (X_test, y_test) = mnist.load_data()
```

```
X_train = X_train.reshape(60000, 784)
X_test = X_test.reshape(10000, 784)
X_train = X_train.astype('float32')
X_test = X_test.astype('float32')
X_train /= 255
X_test /= 255
print(X_train.shape[0], 'train samples')
print(X_test.shape[0], 'test samples')
```

```
num_classes = 10
# convert class vectors to binary class matrices
y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)
```

60000 train samples

10000 test samples

# Fit model

```
model.fit(X_train, y_train, batch_size=128, epochs=10, verbose=1)
```

```
Epoch 1/10  
60000/60000 [=====] - 2s - loss: 0.4996 - acc: 0.8642  
Epoch 2/10  
60000/60000 [=====] - 2s - loss: 0.2432 - acc: 0.9318  
Epoch 3/10  
60000/60000 [=====] - 1s - loss: 0.2004 - acc: 0.9429  
Epoch 4/10  
60000/60000 [=====] - 1s - loss: 0.1734 - acc: 0.9500  
Epoch 5/10  
60000/60000 [=====] - 1s - loss: 0.1540 - acc: 0.9554  
Epoch 6/10  
60000/60000 [=====] - 1s - loss: 0.1393 - acc: 0.9597  
Epoch 7/10  
60000/60000 [=====] - 1s - loss: 0.1273 - acc: 0.9627  
Epoch 8/10  
60000/60000 [=====] - 1s - loss: 0.1172 - acc: 0.9656  
Epoch 9/10  
60000/60000 [=====] - 1s - loss: 0.1088 - acc: 0.9683  
Epoch 10/10  
60000/60000 [=====] - 1s - loss: 0.1020 - acc: 0.9703
```

# Fit with Validation

```
model.fit(X_train, y_train, batch_size=128, epochs=10, verbose=1, validation_split=.1)|
```

```
Train on 54000 samples, validate on 6000 samples
Epoch 1/10
54000/54000 [=====] - 2s - loss: 0.5146 - acc: 0.8616 - val_loss: 0.2425 - val_acc: 0.9322
Epoch 2/10
54000/54000 [=====] - 1s - loss: 0.2618 - acc: 0.9266 - val_loss: 0.1934 - val_acc: 0.9442
Epoch 3/10
54000/54000 [=====] - 1s - loss: 0.2161 - acc: 0.9397 - val_loss: 0.1717 - val_acc: 0.9537
Epoch 4/10
54000/54000 [=====] - 1s - loss: 0.1879 - acc: 0.9470 - val_loss: 0.1519 - val_acc: 0.9570
Epoch 5/10
54000/54000 [=====] - 1s - loss: 0.1676 - acc: 0.9528 - val_loss: 0.1440 - val_acc: 0.9603
Epoch 6/10
54000/54000 [=====] - 1s - loss: 0.1506 - acc: 0.9566 - val_loss: 0.1296 - val_acc: 0.9638
Epoch 7/10
54000/54000 [=====] - 1s - loss: 0.1378 - acc: 0.9603 - val_loss: 0.1281 - val_acc: 0.9627
Epoch 8/10
54000/54000 [=====] - 1s - loss: 0.1268 - acc: 0.9626 - val_loss: 0.1177 - val_acc: 0.9660
Epoch 9/10
54000/54000 [=====] - 1s - loss: 0.1175 - acc: 0.9659 - val_loss: 0.1159 - val_acc: 0.9657
Epoch 10/10
54000/54000 [=====] - 1s - loss: 0.1096 - acc: 0.9677 - val loss: 0.1131 - val acc: 0.9662
```

# Evaluating on Test Set

```
score = model.evaluate(X_test, y_test, verbose=0)
print("Test loss: {:.3f}".format(score[0]))
print("Test Accuracy: {:.3f}".format(score[1]))
```

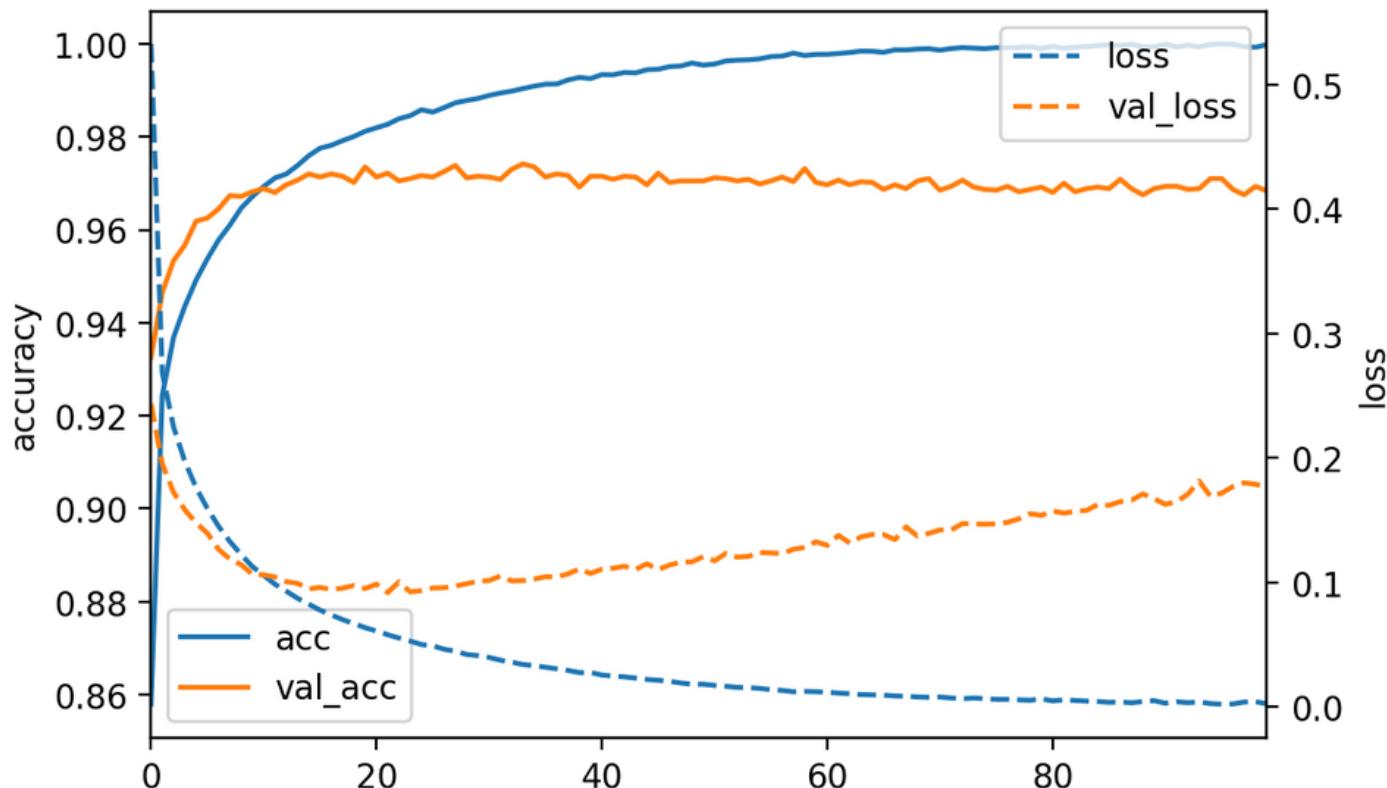
Test loss: 0.120

Test Accuracy: 0.966

# Loggers and Callbacks

```
model = Sequential([
    Dense(32, input_shape=(784,), activation='relu'),
    Dense(10, activation='softmax'),
])
model.compile("adam", "categorical_crossentropy", metrics=['accuracy'])
history_callback = model.fit(X_train, y_train, batch_size=128,
                             epochs=100, verbose=1, validation_split=.1)
```

```
pd.DataFrame(history_callback.history).plot()
```



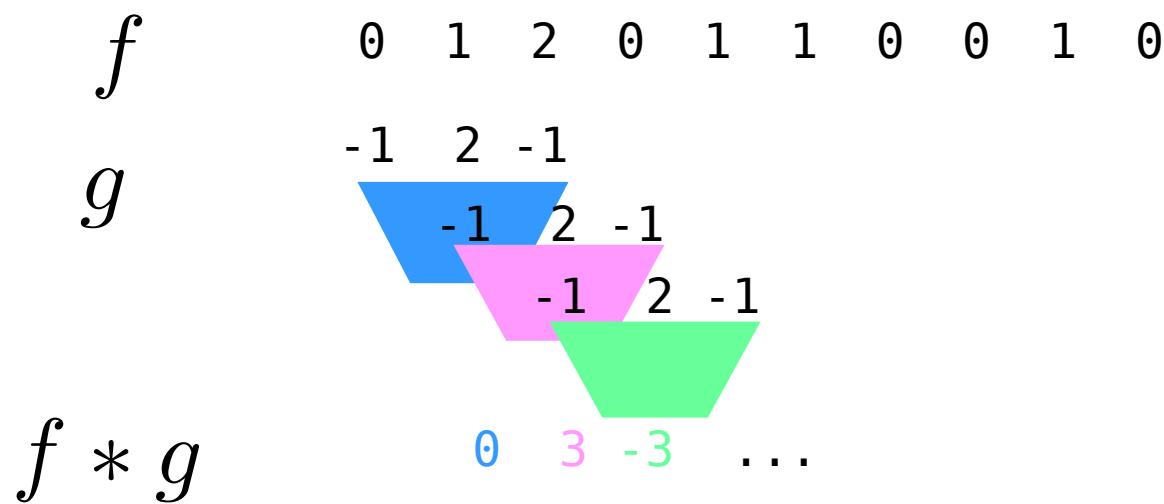
# Convolutional neural networks

# Idea

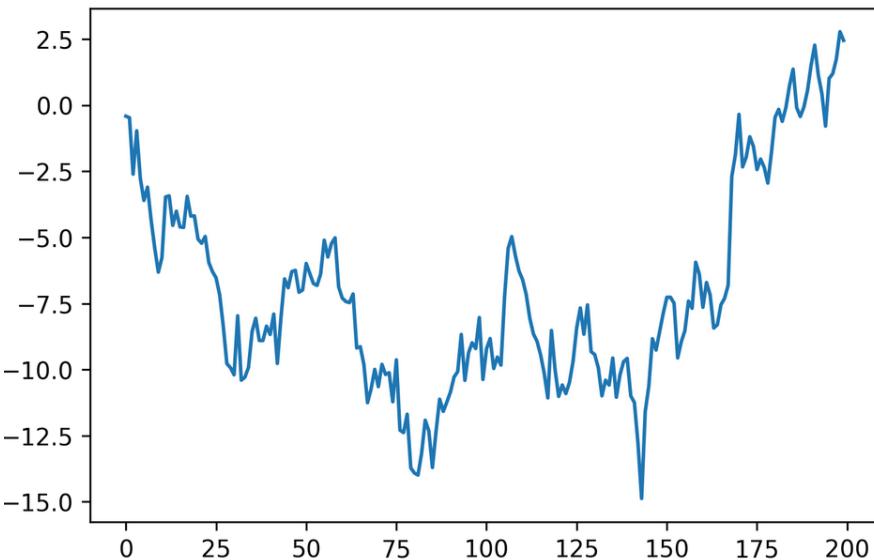
- Translation invariance
- Weight sharing

# Definition of Convolution

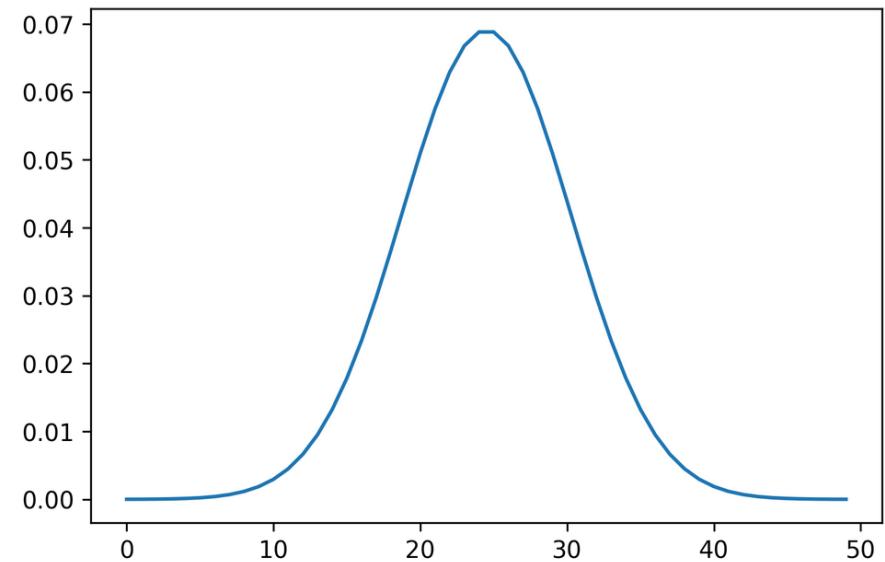
$$\begin{aligned}(f * g)[n] &= \sum_{m=-\infty}^{\infty} f[m] g[n - m] \\&= \sum_{m=-\infty}^{\infty} f[n - m] g[m].\end{aligned}$$



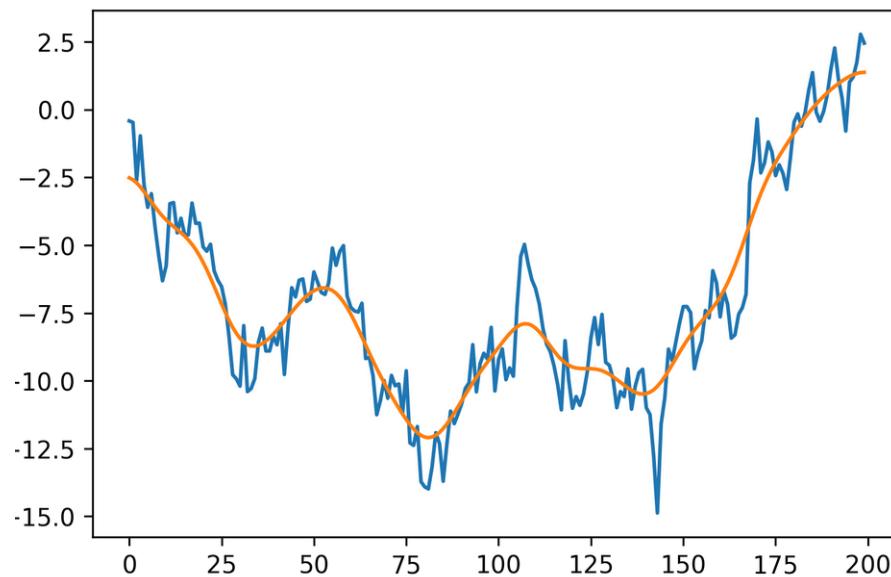
# 1d example: Gaussian smoothing



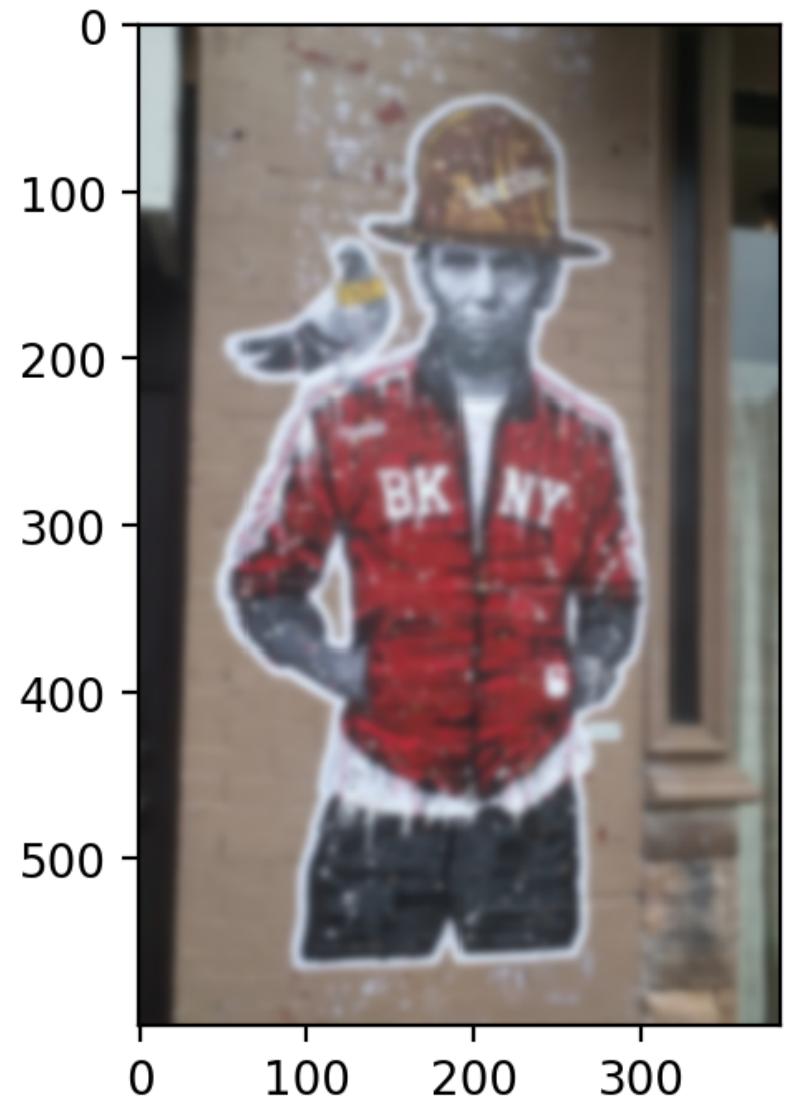
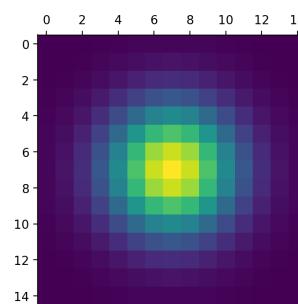
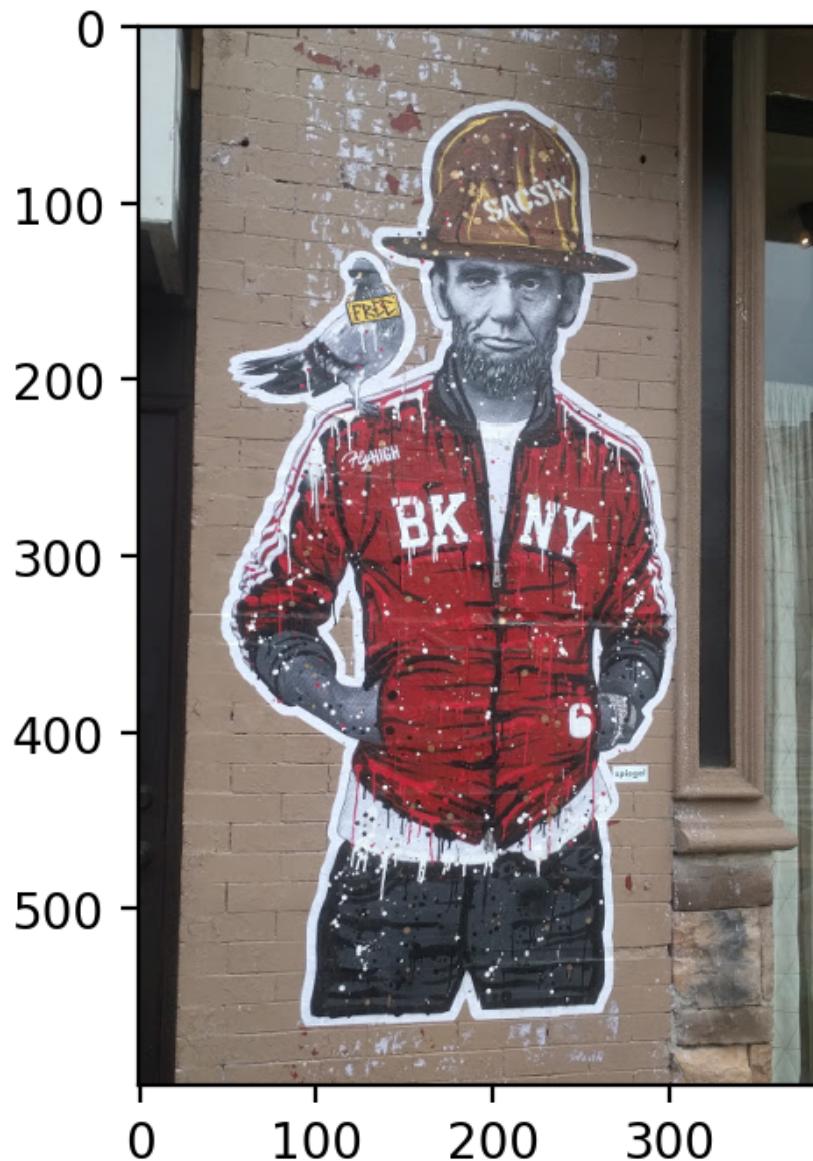
\*



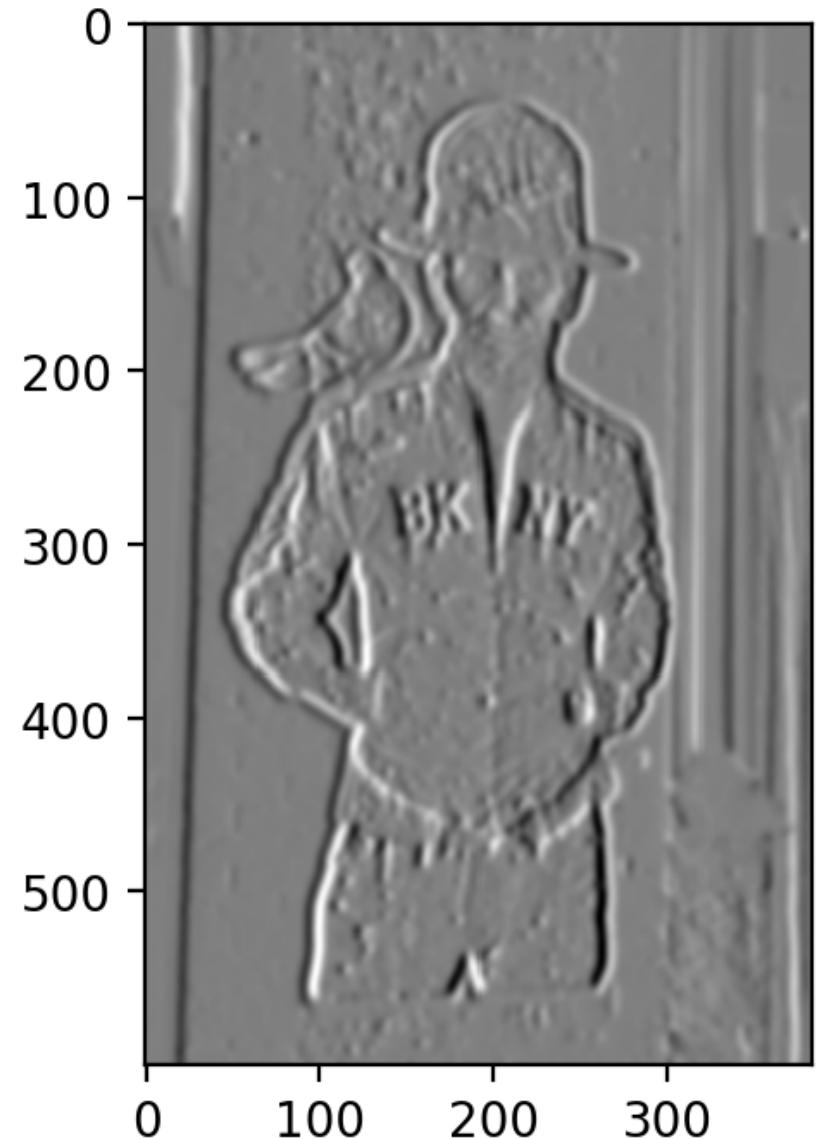
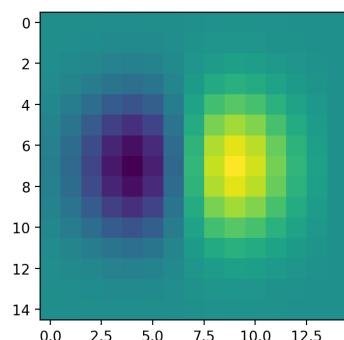
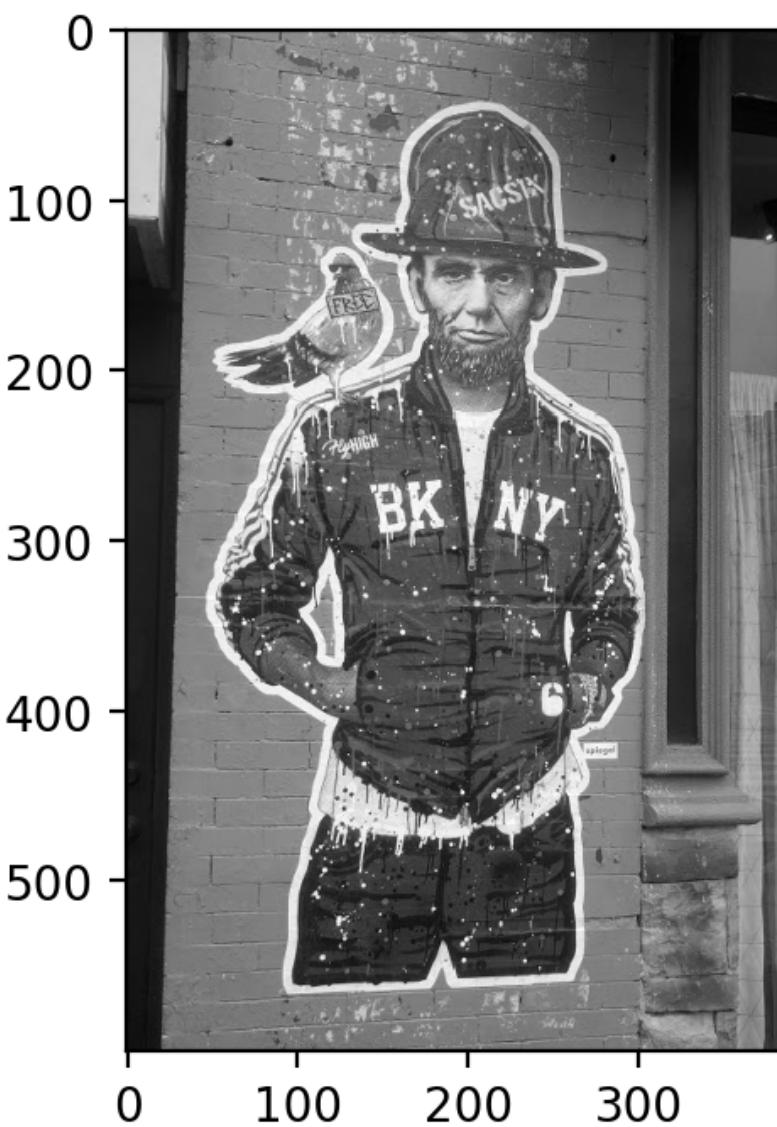
=



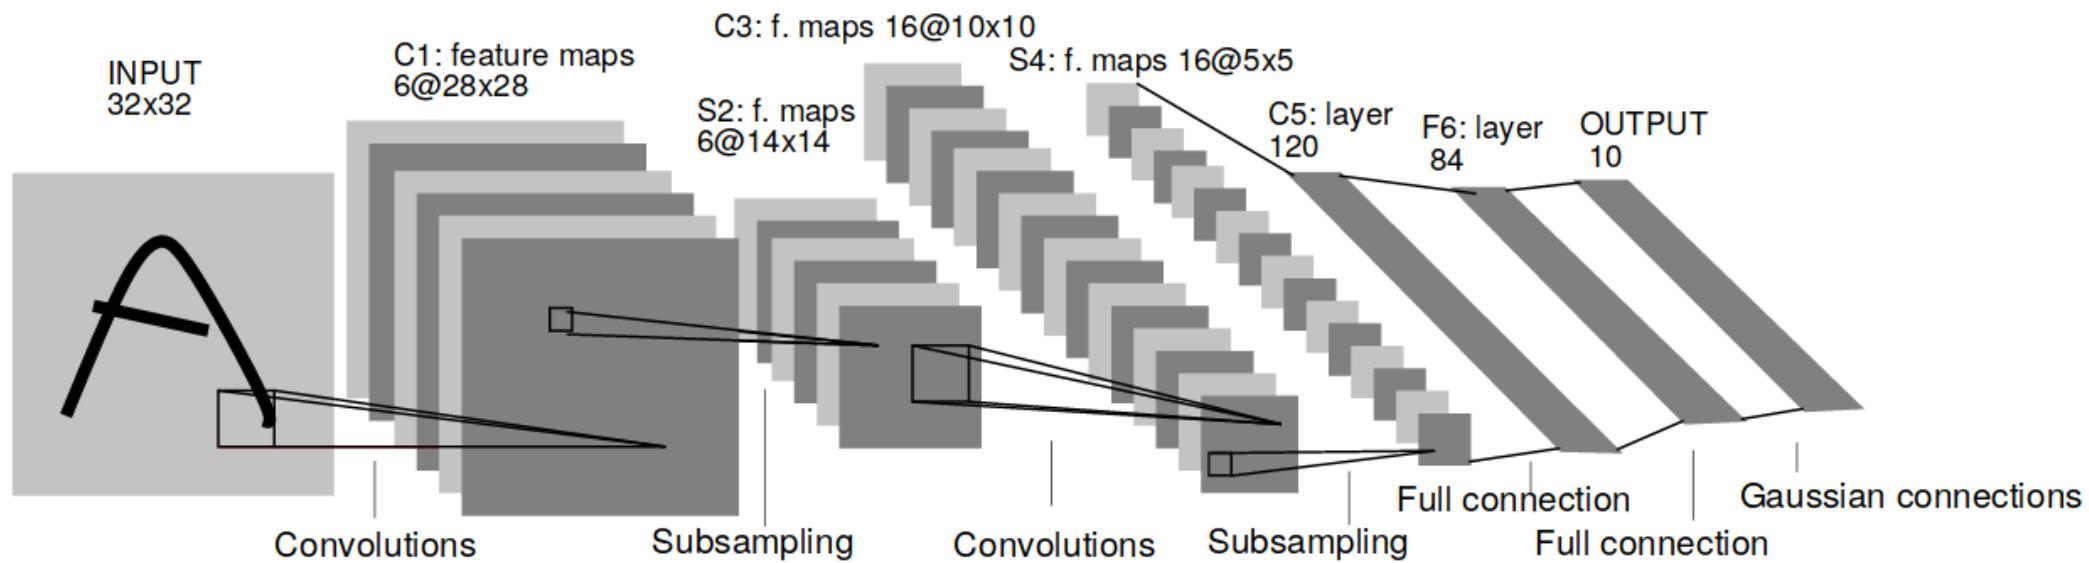
# 2d smoothing



# 2d Gradients

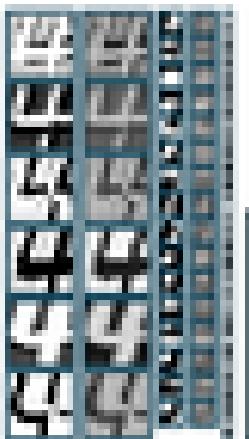


# Convolutional Neural Networks

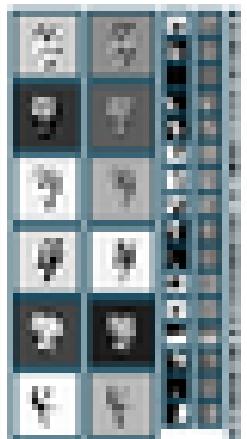
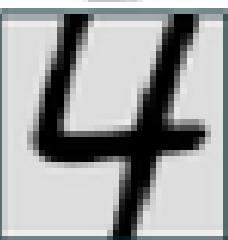


Y. LeCun, L. Bottou, Y. Bengio, and P. Haffner.

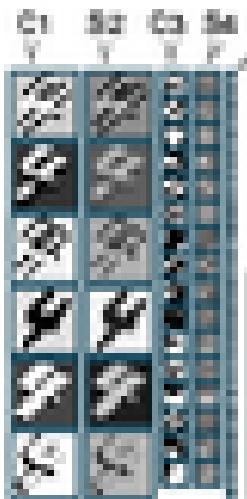
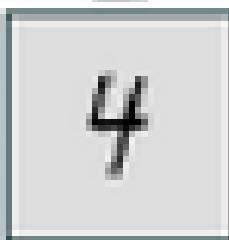
Gradient-based learning applied to document recognition



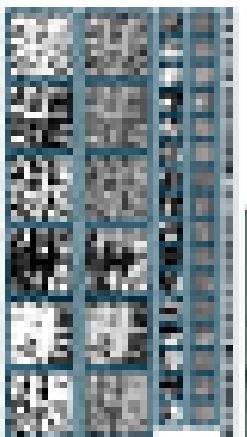
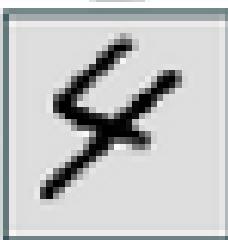
4



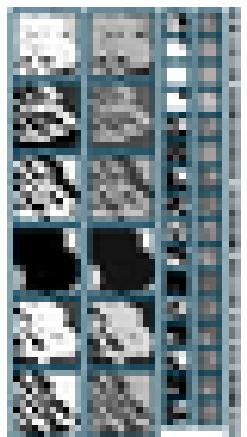
4



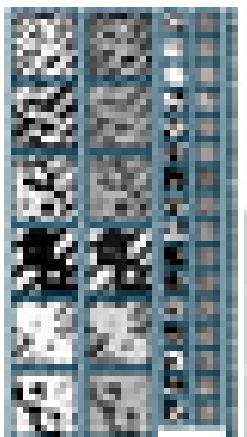
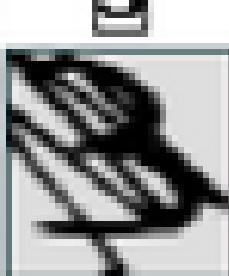
1



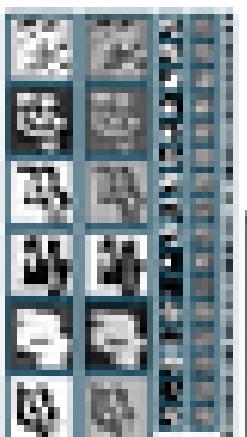
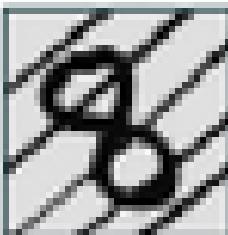
1



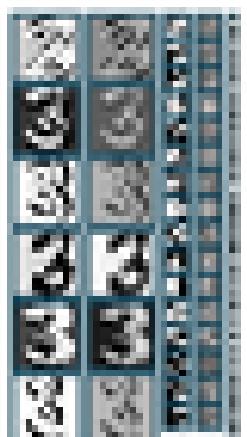
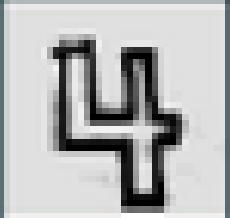
2



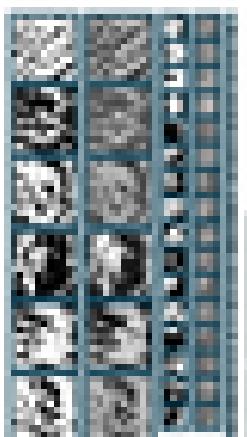
1



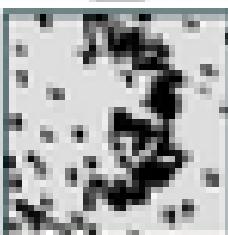
4



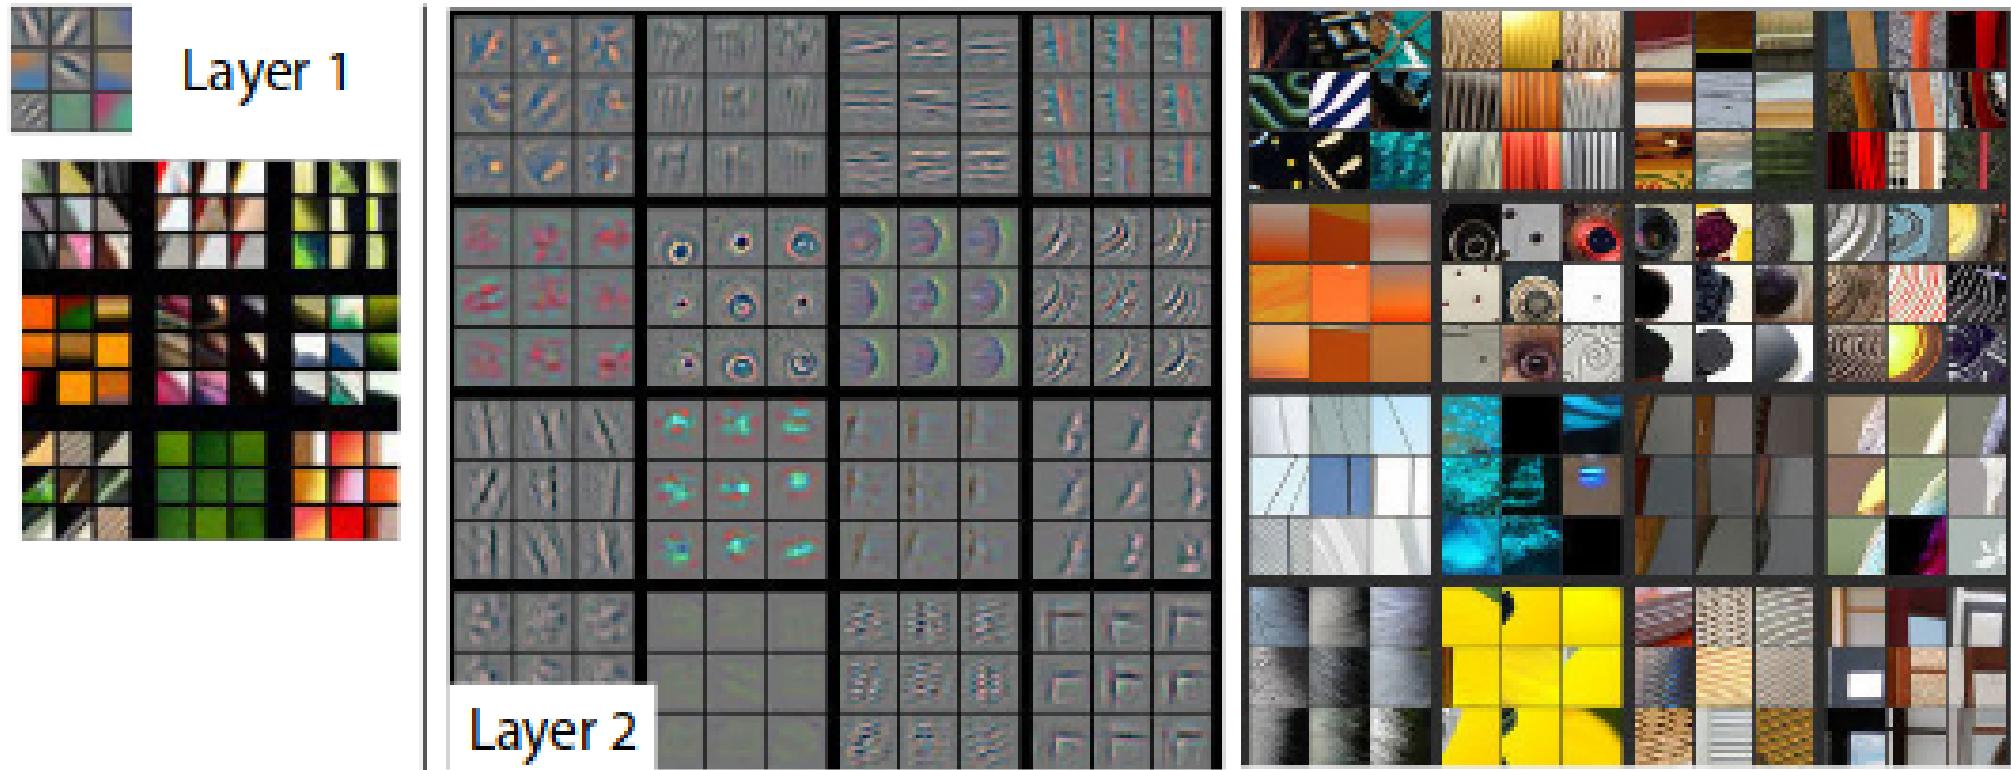
4



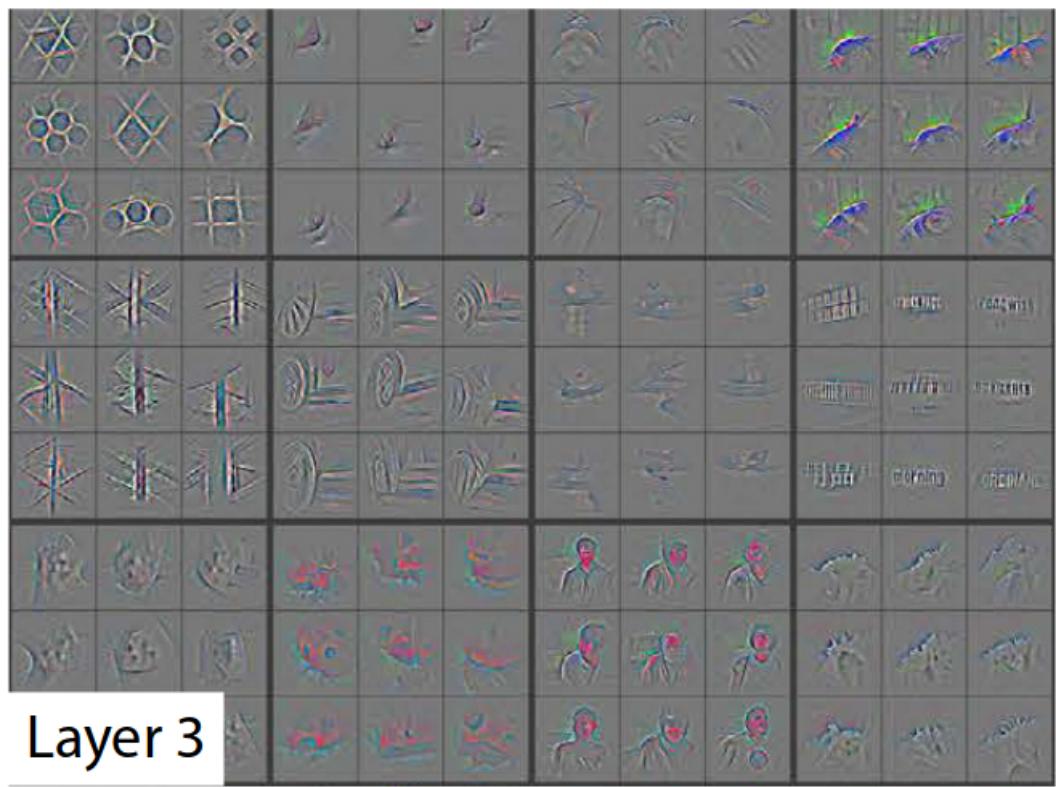
2



# Deconvolution

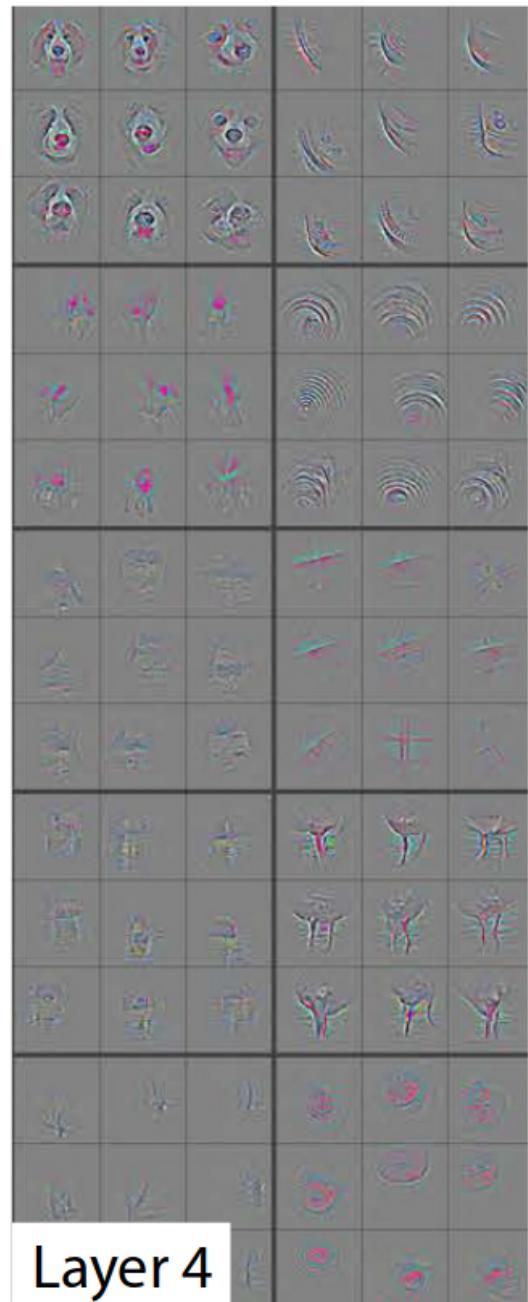


<https://www.cs.nyu.edu/~fergus/papers/zeilerECCV2014.pdf>

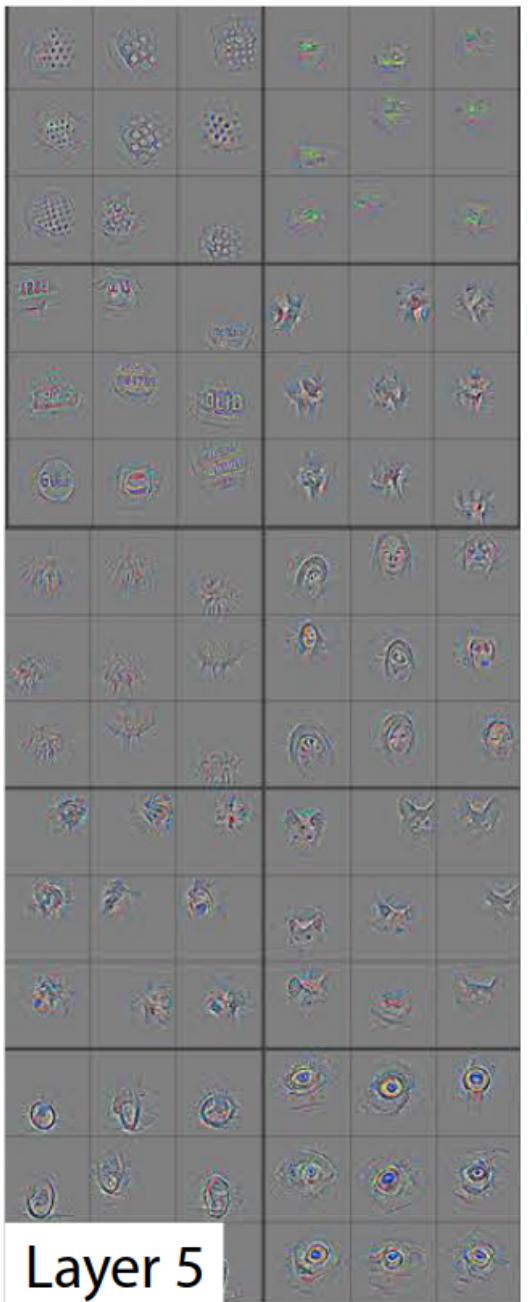


## Layer 3





Layer 4



Layer 5



# Conv-nets with keras

# Preparing data

```
batch_size = 128
num_classes = 10
epochs = 12

# input image dimensions
img_rows, img_cols = 28, 28

# the data, shuffled and split between train and test sets
(x_train, y_train), (x_test, y_test) = mnist.load_data()

x_train_images = x_train.reshape(x_train.shape[0], img_rows, img_cols, 1)
x_test_images = x_test.reshape(x_test.shape[0], img_rows, img_cols, 1)
input_shape = (img_rows, img_cols, 1)

y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)
```

Channels

# Create tiny network

```
from keras.layers import Conv2D, MaxPooling2D, Flatten  
  
num_classes = 10  
cnn_small = Sequential()  
cnn_small.add(Conv2D(8, kernel_size=(3, 3),  
                    activation='relu',  
                    input_shape=input_shape))  
cnn_small.add(MaxPooling2D(pool_size=(2, 2)))  
cnn_small.add(Conv2D(8, (3, 3), activation='relu'))  
cnn_small.add(MaxPooling2D(pool_size=(2, 2)))  
cnn_small.add(Flatten())  
cnn_small.add(Dense(64, activation='relu'))  
cnn_small.add(Dense(num_classes, activation='softmax'))
```

# Number of Parameters

## Convolutional network for MNIST

```
cnn_small.summary()
```

Layer (type)	Output Shape	Param #
conv2d_9 (Conv2D)	(None, 26, 26, 8)	80
max_pooling2d_9 (MaxPooling2D)	(None, 13, 13, 8)	0
conv2d_10 (Conv2D)	(None, 11, 11, 8)	584
max_pooling2d_10 (MaxPooling2D)	(None, 5, 5, 8)	0
flatten_5 (Flatten)	(None, 200)	0
dense_172 (Dense)	(None, 64)	12864
dense_173 (Dense)	(None, 10)	650

Total params: 14,178.0  
Trainable params: 14,178.0  
Non-trainable params: 0.0

## Dense network for MNIST

Layer (type)	Output Shape	Param #
dense_1 (Dense)	(None, 512)	401920
dropout_1 (Dropout)	(None, 512)	0
dense_2 (Dense)	(None, 512)	262656
dropout_2 (Dropout)	(None, 512)	0
dense_3 (Dense)	(None, 10)	5130

Total params: 669,706.0  
Trainable params: 669,706.0  
Non-trainable params: 0.0

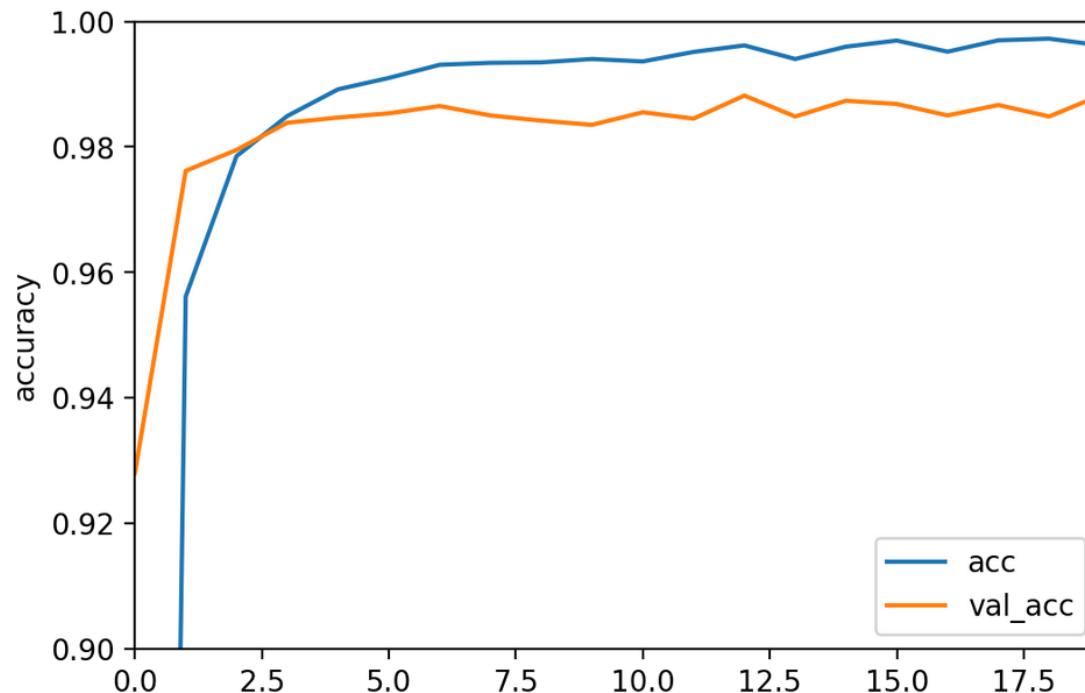
# Train and Evaluate

```
cnn.compile("adam", "categorical_crossentropy", metrics=['accuracy'])
history_cnn = cnn.fit(X_train_images, y_train,
                       batch_size=128, epochs=20, verbose=1, validation_split=.1)
```

```
: cnn.evaluate(X_test_images, y_test)
```

```
9952/10000 [=====>.] - ETA: 0s
```

```
[0.089020583277629253, 0.9842999999999995]
```

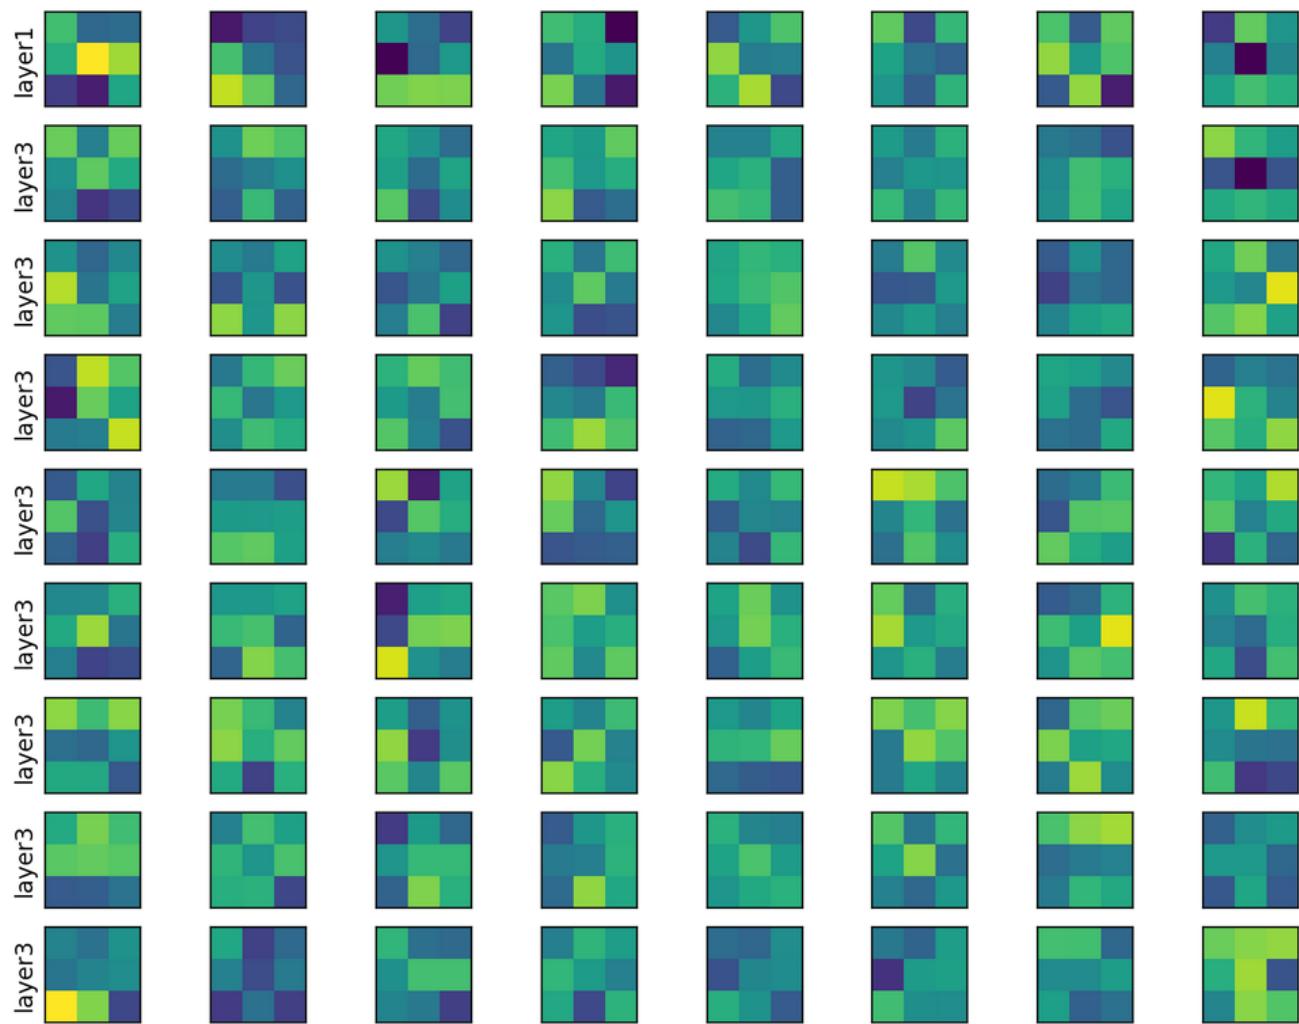


# Visualize filters

```
weights, biases = cnn_small.layers[0].get_weights()  
weights2, biases2 = cnn_small.layers[2].get_weights()  
print(weights.shape)  
print(weights2.shape)
```

(3, 3, 1, 8)

(3, 3, 8, 8)





# Batch Normalization

**Input:** Values of  $x$  over a mini-batch:  $\mathcal{B} = \{x_1 \dots m\}$ ;

Parameters to be learned:  $\gamma, \beta$

**Output:**  $\{y_i = \text{BN}_{\gamma, \beta}(x_i)\}$

$$\mu_{\mathcal{B}} \leftarrow \frac{1}{m} \sum_{i=1}^m x_i \quad // \text{mini-batch mean}$$

$$\sigma_{\mathcal{B}}^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2 \quad // \text{mini-batch variance}$$

$$\hat{x}_i \leftarrow \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}} \quad // \text{normalize}$$

$$y_i \leftarrow \gamma \hat{x}_i + \beta \equiv \text{BN}_{\gamma, \beta}(x_i) \quad // \text{scale and shift}$$

# VGG and Imagenet Filters

# Inspecting VGG

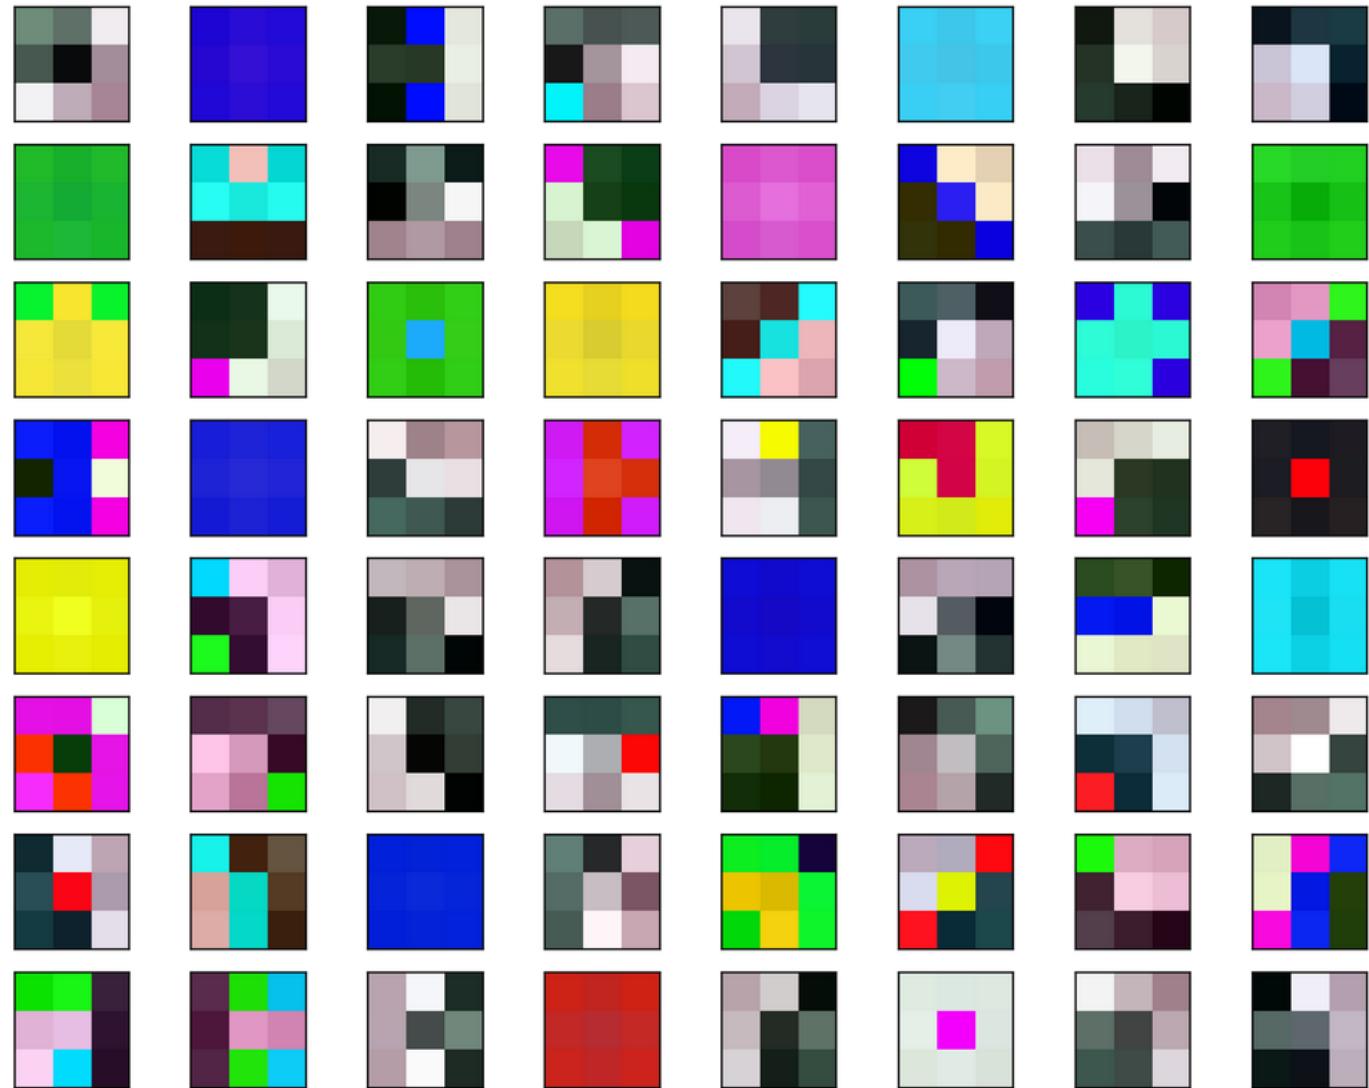
```
from keras import applications  
  
# build the VGG16 network  
model = applications.VGG16(include_top=False,  
                           weights='imagenet')  
  
model.summary()
```

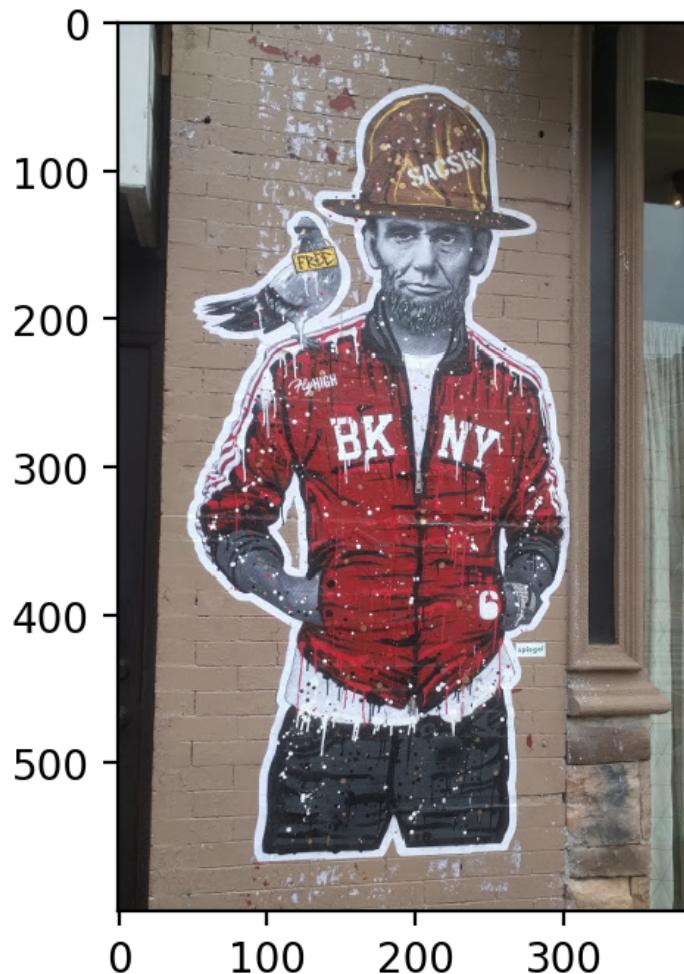
Layer (type)	Output Shape	Param #
input_1 (InputLayer)	(None, None, None, 3)	0
block1_conv1 (Conv2D)	(None, None, None, 64)	1792
block1_conv2 (Conv2D)	(None, None, None, 64)	36928
block1_pool (MaxPooling2D)	(None, None, None, 64)	0
block2_conv1 (Conv2D)	(None, None, None, 128)	73856
block2_conv2 (Conv2D)	(None, None, None, 128)	147584
block2_pool (MaxPooling2D)	(None, None, None, 128)	0
block3_conv1 (Conv2D)	(None, None, None, 256)	295168
block3_conv2 (Conv2D)	(None, None, None, 256)	590080
block3_conv3 (Conv2D)	(None, None, None, 256)	590080
block3_pool (MaxPooling2D)	(None, None, None, 256)	0
block4_conv1 (Conv2D)	(None, None, None, 512)	1180160
block4_conv2 (Conv2D)	(None, None, None, 512)	2359808
block4_conv3 (Conv2D)	(None, None, None, 512)	2359808
block4_pool (MaxPooling2D)	(None, None, None, 512)	0
block5_conv1 (Conv2D)	(None, None, None, 512)	2359808
block5_conv2 (Conv2D)	(None, None, None, 512)	2359808
block5_conv3 (Conv2D)	(None, None, None, 512)	2359808
block5_pool (MaxPooling2D)	(None, None, None, 512)	0
<hr/>		
Total params: 14,714,688.0		
Trainable params: 14,714,688.0		
Non-trainable params: 0.0		

# VGG filters

```
vgg_weights, vgg_biases = model.layers[1].get_weights()  
vgg_weights.shape
```

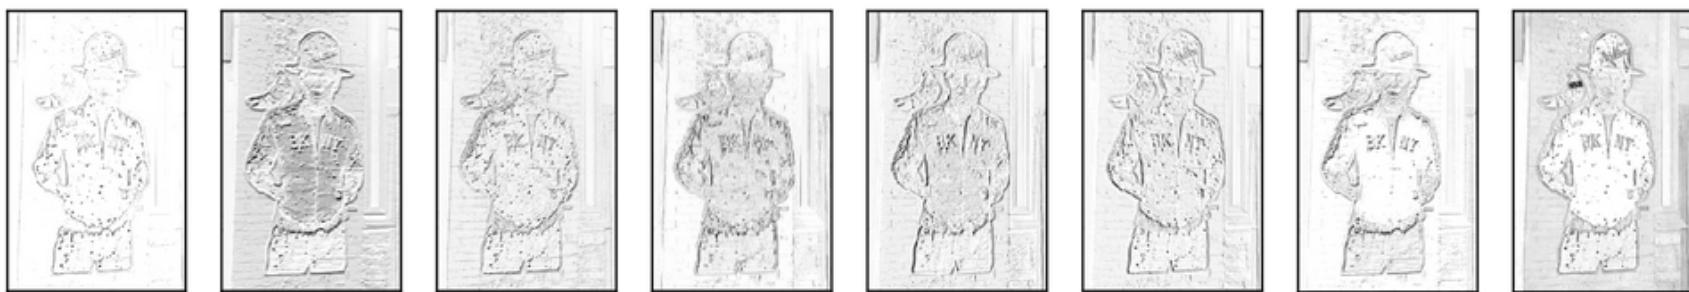
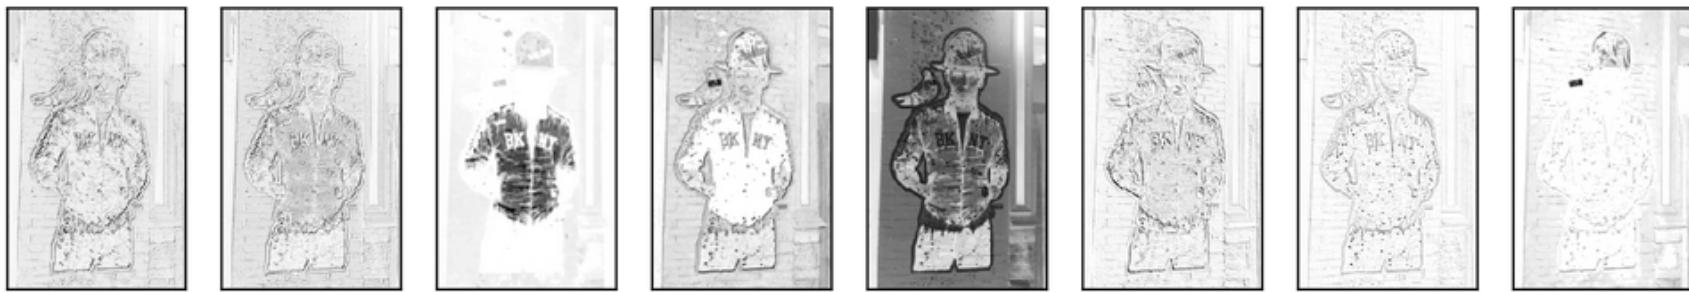
(3, 3, 3, 64)





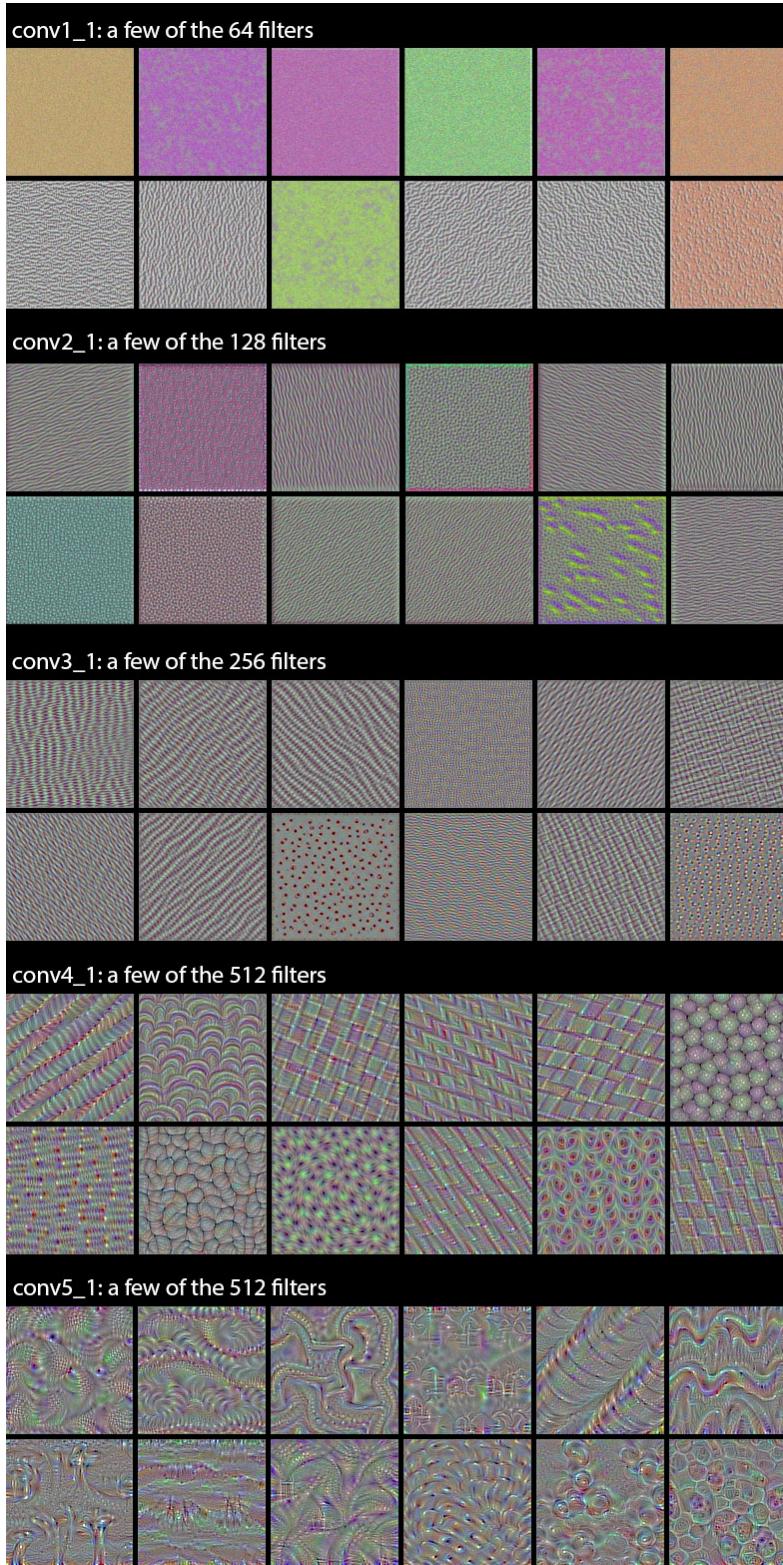
```
get_3rd_layer_output = K.function([model.layers[0].input],  
                                 [model.layers[3].output])  
get_6rd_layer_output = K.function([model.layers[0].input],  
                                 [model.layers[6].output])  
  
layer3_output = get_3rd_layer_output([[image]])[0]  
layer6_output = get_6rd_layer_output([[image]])[0]  
  
print(layer3_output.shape)  
print(layer6_output.shape)  
  
(1, 300, 192, 64)  
(1, 150, 96, 128)
```

after first pooling layer



after second pooling layer

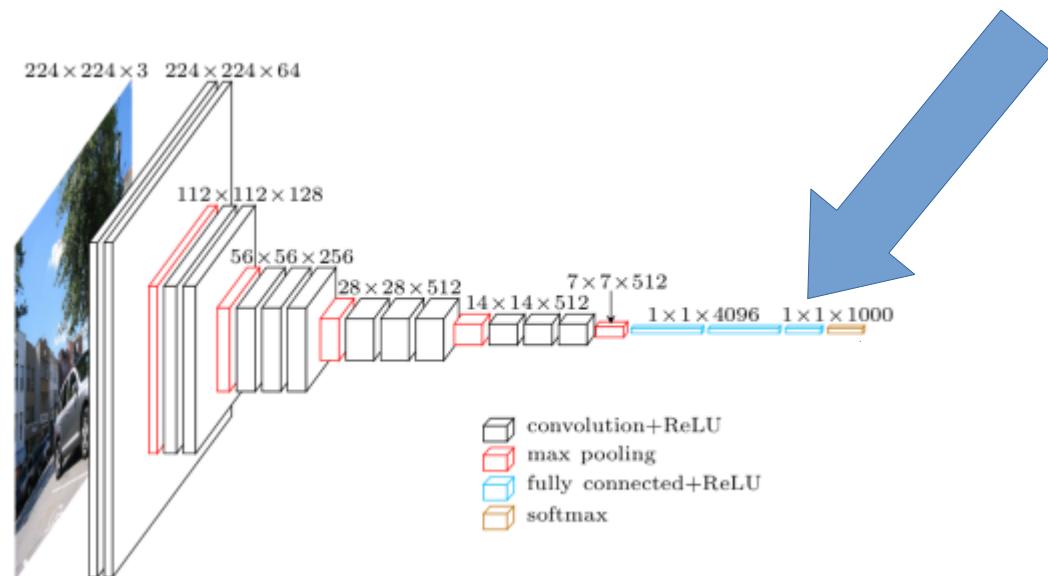




<https://blog.keras.io/how-convolutional-neural-networks-see-the-world.html>

# Using pre-trained networks

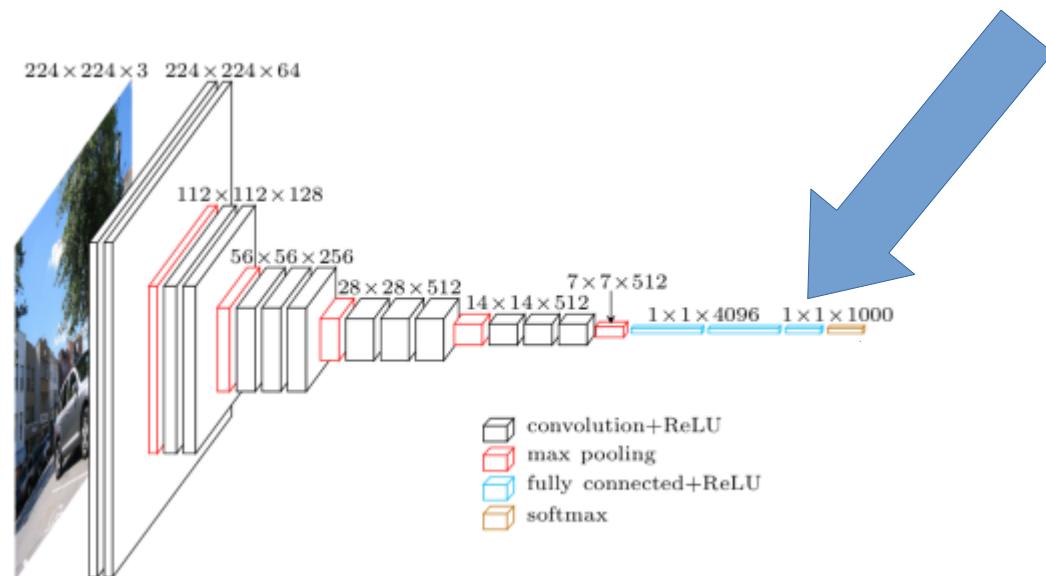
- Train on “large enough” data.
- Apply to new “small” dataset.
- Take activations of last or second to last fully connected layer.



See <http://cs231n.github.io/transfer-learning/>

# Using pre-trained networks

- Train on “large enough” data.
- Apply to new “small” dataset.
- Take activations of last or second to last fully connected layer.



# Ball snake vs Carpet Python



```
import flickrapi
flickr = flickrapi.FlickrAPI(api_key, api_secret, format='json')

def search_ids(search_string="python", per_page=10):
    photos_response = flickr.photos.search(text=search_string, per_page=per_page, sort='relevance')
    photos = json.loads(photos_response.decode('utf-8'))['photos']['photo']
    ids = [photo['id'] for photo in photos]
    return ids

def get_url(photo_id="33510015330"):
    response = flickr.photos.getsizes(photo_id=photo_id)
    sizes = json.loads(response.decode('utf-8'))['sizes']['size']
    for size in sizes:
        if size['label'] == "Small":
            return size['source']

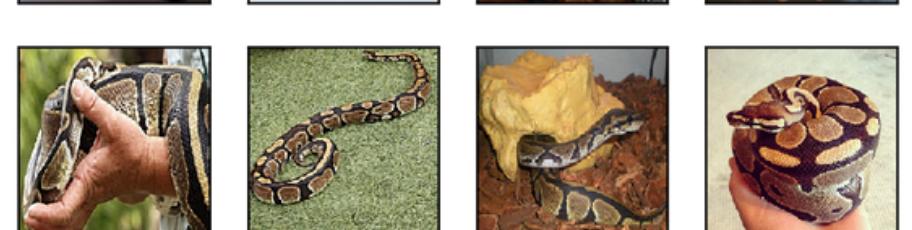
ids = search_ids("ball snake", per_page=100)
urls_ball = [get_url(photo_id=i) for i in ids]

from urllib.request import urlretrieve
import os
for url in urls_carpet:
    urlretrieve(url, os.path.join("snakes", "carpet", os.path.basename(url)))
```

Carpet Python (100 total)



Ball Snake (100 total)



# Extracting Features using VGG

```
from keras.preprocessing import image

images_carpet = [image.load_img(os.path.join("snakes", "carpet", os.path.basename(url)), target_size=(224, 224))
                 for url in urls_carpet]
images_ball = [image.load_img(os.path.join("snakes", "ball", os.path.basename(url)), target_size=(224, 224))
                 for url in urls_ball]
X = np.array([image.img_to_array(img) for img in images_carpet + images_ball])
```

```
model = applications.VGG16(include_top=False,
                            weights='imagenet')
```

```
X.shape
```

```
(200, 224, 224, 3)
```

```
from keras.applications.vgg16 import preprocess_input
X_pre = preprocess_input(X)
features = model.predict(X_pre)
```

```
features.shape
```

```
(200, 7, 7, 512)
```

```
features_ = features.reshape(200, -1)
```

# Classification with Logreg

```
from sklearn.model_selection import train_test_split
y = np.zeros(200, dtype='int')
y[100:] = 1
X_train, X_test, y_train, y_test = train_test_split(features_, y, stratify=y)
```

```
from sklearn.linear_model import LogisticRegressionCV
lr = LogisticRegressionCV().fit(X_train, y_train)
```

```
print(lr.score(X_train, y_train))
```

```
1.0
```

```
print(lr.score(X_test, y_test))
```

```
0.82
```

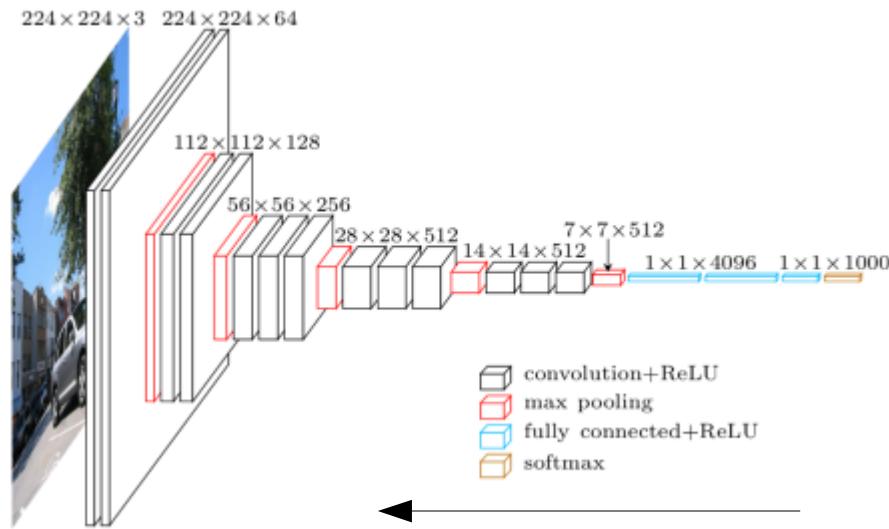
```
from sklearn.metrics import confusion_matrix
confusion_matrix(y_test, lr.predict(X_test))
```

```
array([[24,  1],
       [ 8, 17]])
```

Caveat: I haven't really checked if these make sense.

# Finetuning

- Start with pre-trained net
- Back-propagate error through all layers
- “tune” filters to new data.



See <https://keras.io/applications/>